

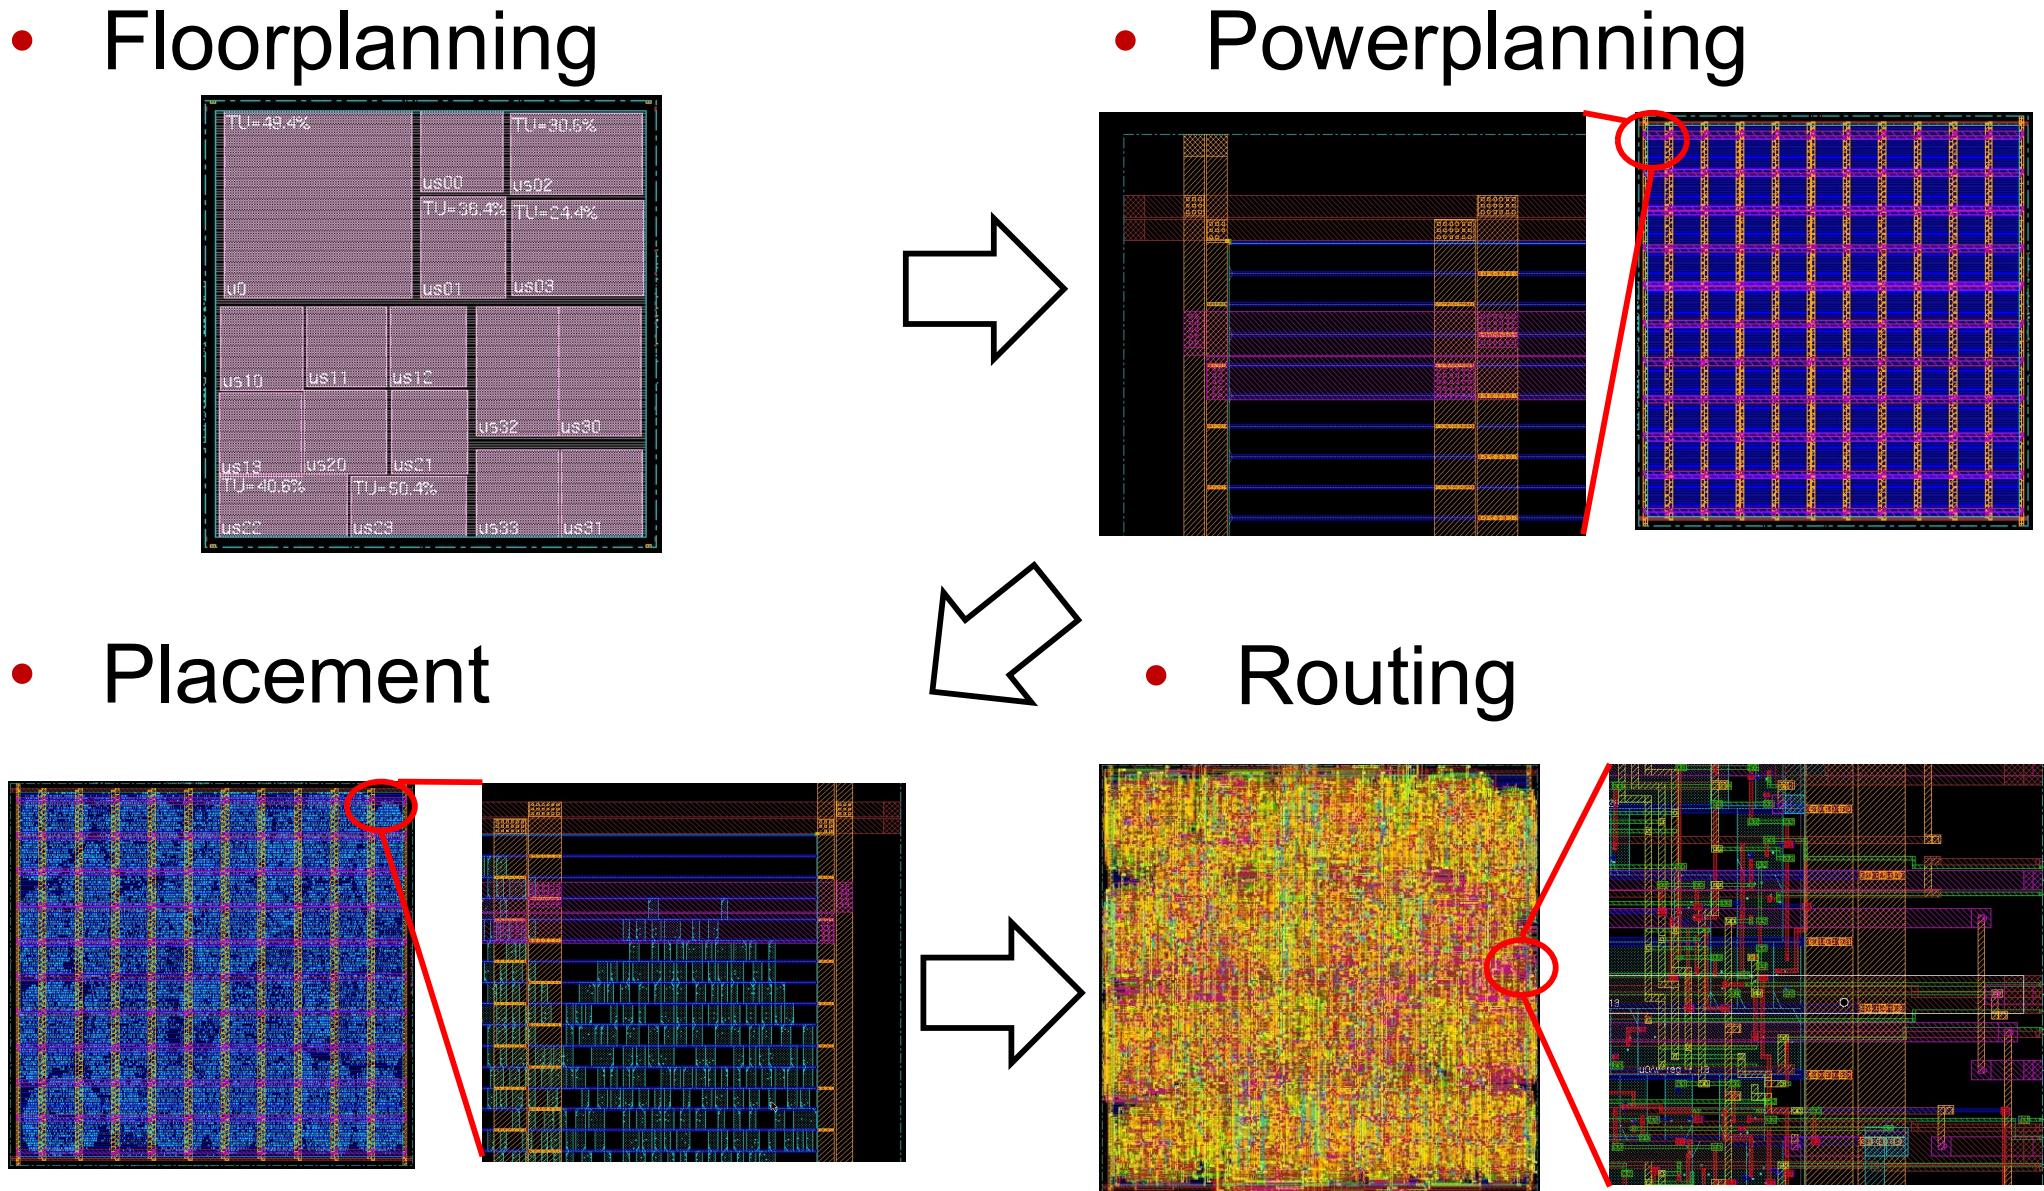
ECE 260C, Spring 2025

Floorplan and PDN

Andrew B. Kahng

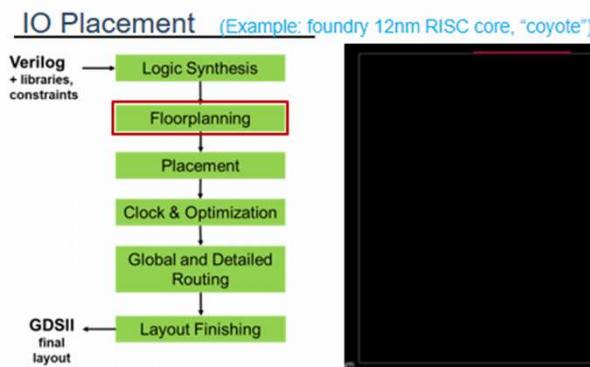
Thanks to: Zhiang Wang, Peter Gadfort

Physical Design Flow Pictures (old ECE 260B slide)

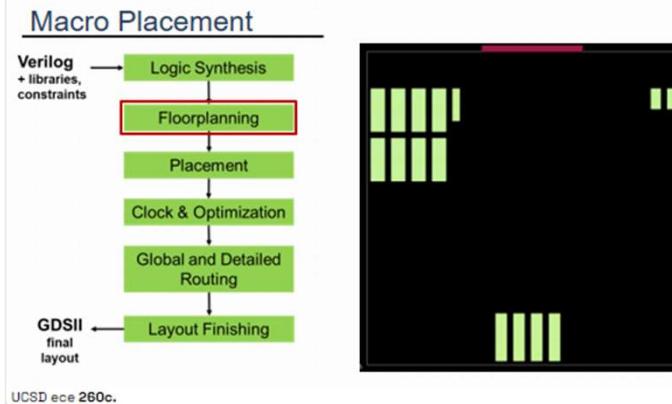


Floorplanning

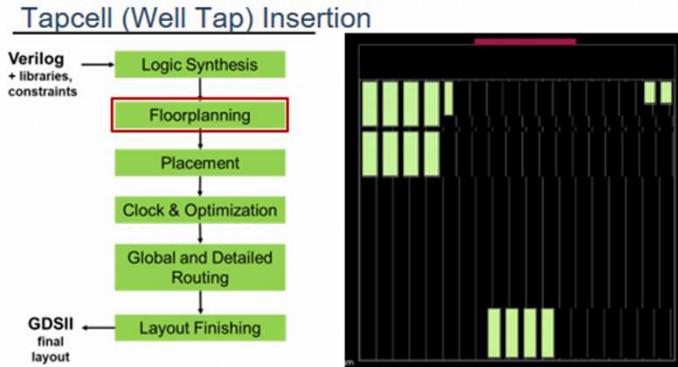
Recall: “Floorplanning” Performs Many Tasks !



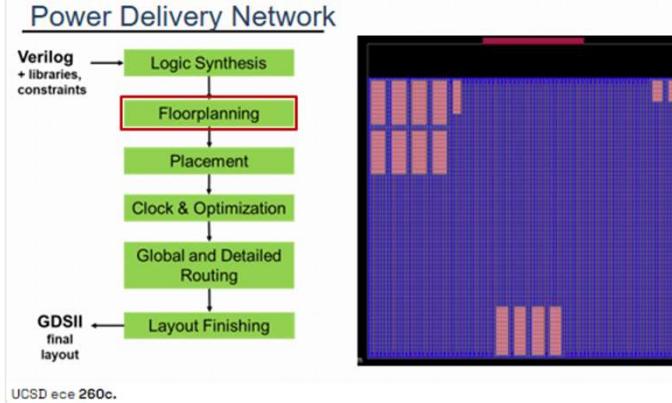
UCSD ece 260c.



UCSD ece 260c.

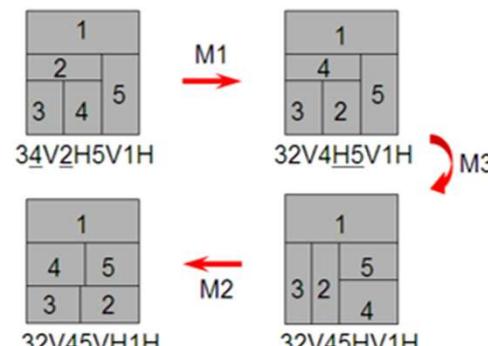


UCSD ece 260c.



UCSD ece 260c.

- <https://openroad.readthedocs.io/en/latest/main/src/ifp/README.html> ++
- IO placement = ppl
- Macro placement = mpl
 - Objectives? Constraints?
 - **Representation? (Realization?) Optimization?**
- Tapcell insertion = tap
- PDN generation = pdn



Fast and Scalable I/O Pin Assignment with Divide-and-Conquer and Hungarian Matching

Publisher: IEEE

Cite This



Vitor Bandeira ; Mateus Fogaça ; Eder Matheus Monteiro ; Isadora Oliveira ; Mingyu Woo ; Ricardo Reis

All /

Pin Placement (ppo)

Fast and Scalable I/O Pin Assignment with Divide-and-Conquer and Hungarian Matching

Vitor Bandeira^{1,3}, Mateus Fogaça^{1,3}, Eder Matheus Monteiro³, Isadora Oliveira^{1,3}, Mingyu Woo⁴ and Ricardo Reis^{1,2,3}

¹PGMicro/²PPGC, ³Instituto de Informática, Universidade Federal do Rio Grande do Sul

⁴ECE Department, UC San Diego, La Jolla, CA, USA

{vvbandeira, mpfogaca, emrmonterio, isoliveira, reis}@inf.ufrgs.br, mwoo@eng.ucsd.edu

Abstract—I/O pin assignment is a crucial task in the floorplanning stage of IC implementation. However, this task has not received much attention in the physical design automation literature. Notably, floorplanning is a highly manual stage of the design flow. Nevertheless, it is known that the impact of I/O pin assignment in total routed wirelength (WL) is in the order of 5%. In advanced nodes, density, power and timing become very crucial and WL impacts of 5% are highly significant. We are therefore motivated to revisit the I/O pin assignment problem in this work. We present a fast and scalable Hungarian matching-based heuristic for I/O pin assignment. We present background scalability studies and a divide-and-conquer strategy that significantly reduces runtime without harm to the quality of results. Our algorithm converges in fewer iterations than previous works and presents superior performance according to criteria from the literature.

pins are implemented using vertical/horizontal metal layers whose pitches are equal to $0.2\mu m$. Then, a core with a perimeter of $6mm$ would have 30,000 candidate locations. Many existing heuristics will face runtime/scalability issues with such high number of candidate locations. Thus, heuristic methods can benefit from criteria to filter candidate locations as well as the divide-and-conquer paradigm.

Our contributions are summarized as follows.

- 1) We devise a scalable Hungarian matching-based [6] approach to I/O pin assignment. Motivated by background scalability studies, we propose a divide-and-conquer strategy to reduce runtime. Our divide-and-conquer strategy can perform I/O pin assignment for a design with 2K I/O pins and 14K candidate positions in less than one second.
- 2) We integrate our I/O pin assigner tool with a global placement

<https://ieeexplore.ieee.org/document/9159791>

ppo = pin placement optimizer

- Where should pins be placed?
 - Chicken-egg: core determines periphery, and vice-versa!
 - Basic approach: “iterate until convergence ...”
 - Sometimes, designer wants to provide guidance/constraints
 - Sometimes, the guidance/constraints are wrong (!) e.g., infeasible

The screenshot shows a web browser displaying the OpenROAD documentation for the Pin Placer module. The URL in the address bar is openroad.readthedocs.io/en/latest/main/src/ppl/README.html. The page title is "Pin Placer". The main content area describes the Pin Placer's function: "Place pins on the boundary of the die on the track grid to minimize net wirelengths. Pin placement also creates a metal shape for each pin using min-area rules." It also notes that for designs with unplaced cells, the net wirelength is computed considering the center of the die area as the unplaced cells position.

The left sidebar contains a navigation menu for the OpenROAD documentation, including sections like "Getting Started with OpenROAD", "Building OpenROAD", "Getting Started with the OpenROAD Flow - OpenROAD-flow-scripts", "Tutorials", "Git Quickstart", "Man pages", "OpenROAD User Guide", "OpenROAD APIs", "Database", "GUI", "Partition Management", "Restructure", "Floorplan Initialization", and "Pin Placement".

The right sidebar lists various commands under the "Commands" section, such as "Define Pin Shape Pattern", "Options", "Face-to-Face direct-bonding IOs", "Set IO Pin Constraints", "Exclude IO Pin Region", "Clear IO Pin Constraints", "Set Pin Length", "Set Pin Length Extension", "Set Pin Thickness Multiplier", "Set Simulated Annealing", "Simulated Annealing Debug Mode", "Place specific Pin", "Place all Pins", "Write Pin Placement", "Useful Developer Commands", "Example scripts", "Regression tests", "Limitations", "References", "FAQs", and "License".

A note box in the main content area provides information about command parameters:

- Parameters in square brackets `[-param param]` are optional.
- Parameters without square brackets `-param2 param2` are required.

The "Define Pin Shape Pattern" section explains that the `define_pin_shape_pattern` command defines a pin placement grid on the specified layer. This grid has positions inside the die area, not only at the edges of the die boundary. A code snippet for the command is shown in a code editor window:

```
define_pin_shape_pattern [-layer layer]
```

Macro Placement: Hier-RTLMP (mpl)

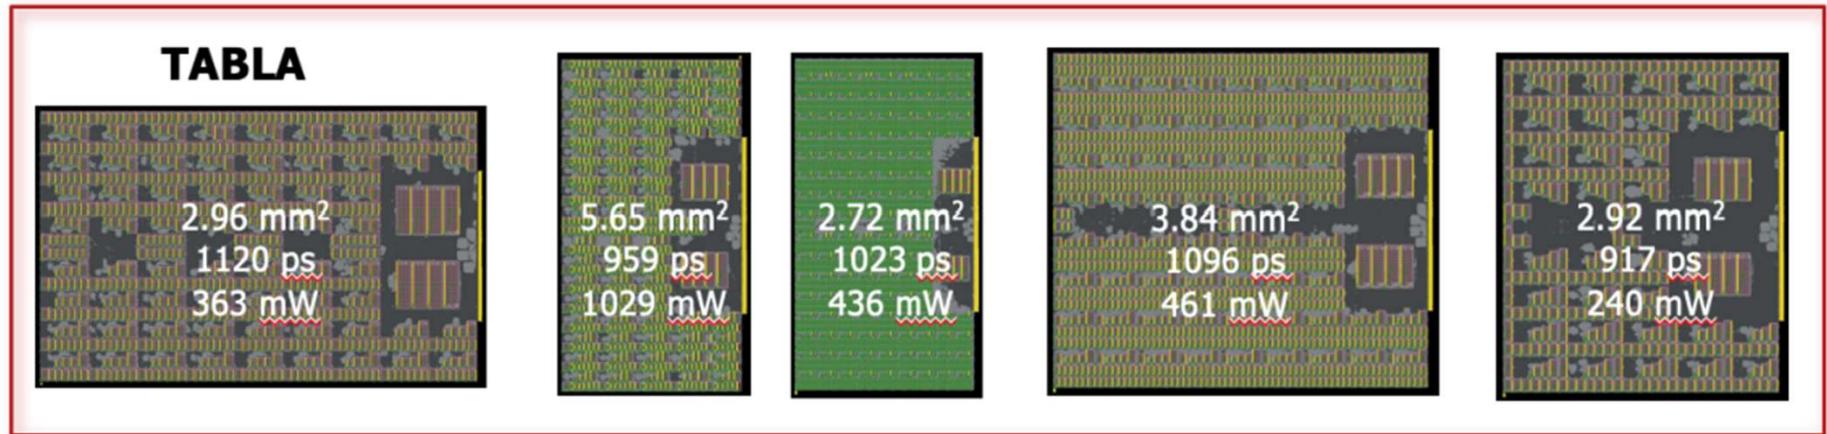
Thanks to Dr. Zhiang Wang zhw033@ucsd.edu

Motivation

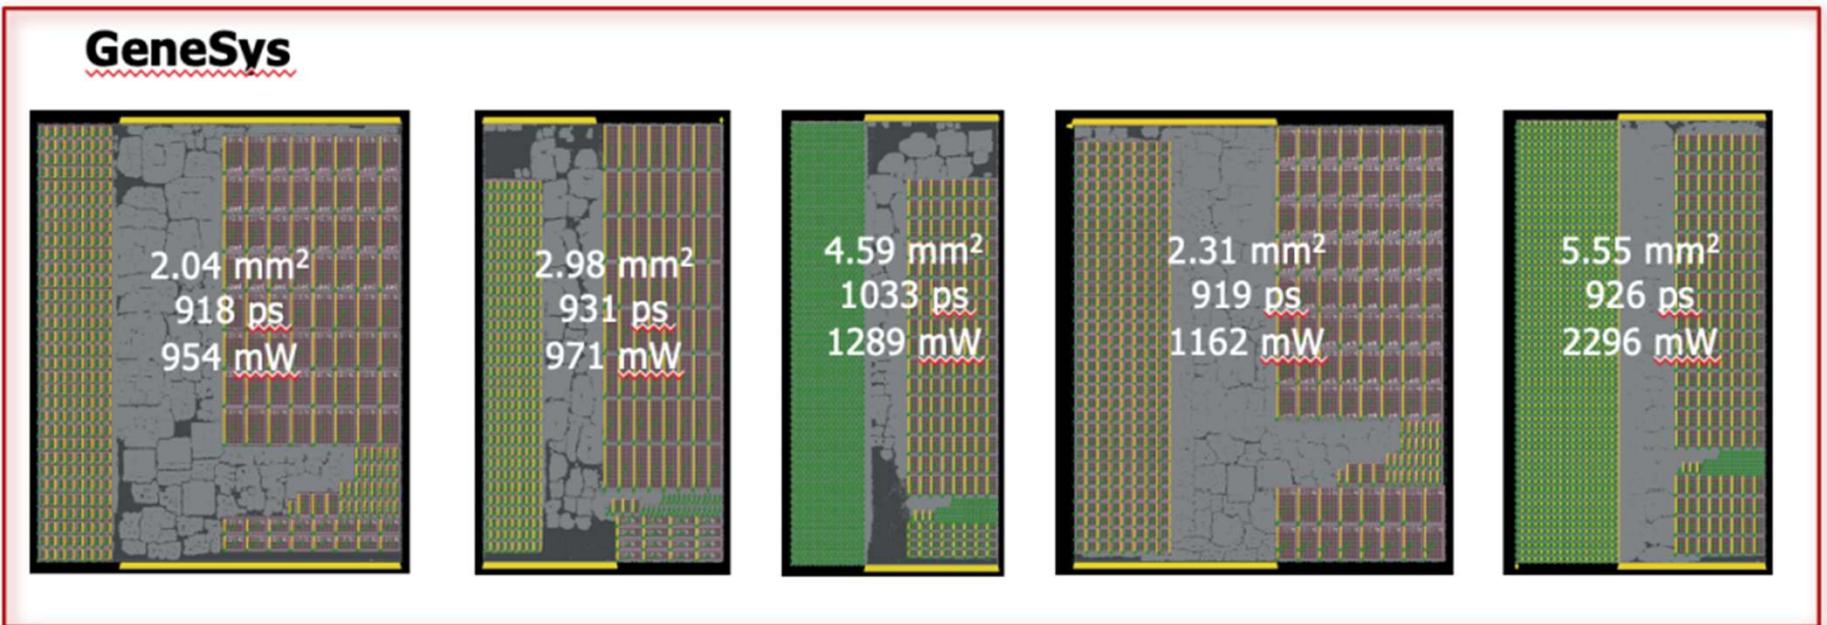
- Macro placement is a critical component of the “handoff” from frontend to the backend of a physical synthesis flow.
- Human experts usually need to consider multiple factors when they do macro placement:
 - Dataflow
 - Logical hierarchy
 - Connectivity between macros and input-output (IO) pins
 - ...
- ML hardware accelerators make the macro placement become more challengeable:
 - Hundreds of macros (examples will be given in the next slide)
 - Complex RTL structure with long, inscrutable autogenerated module names

Macro-Dominated ML Accelerators ([paper: c397](#))

TABLA: machine learning accelerators for non-DNN algorithms

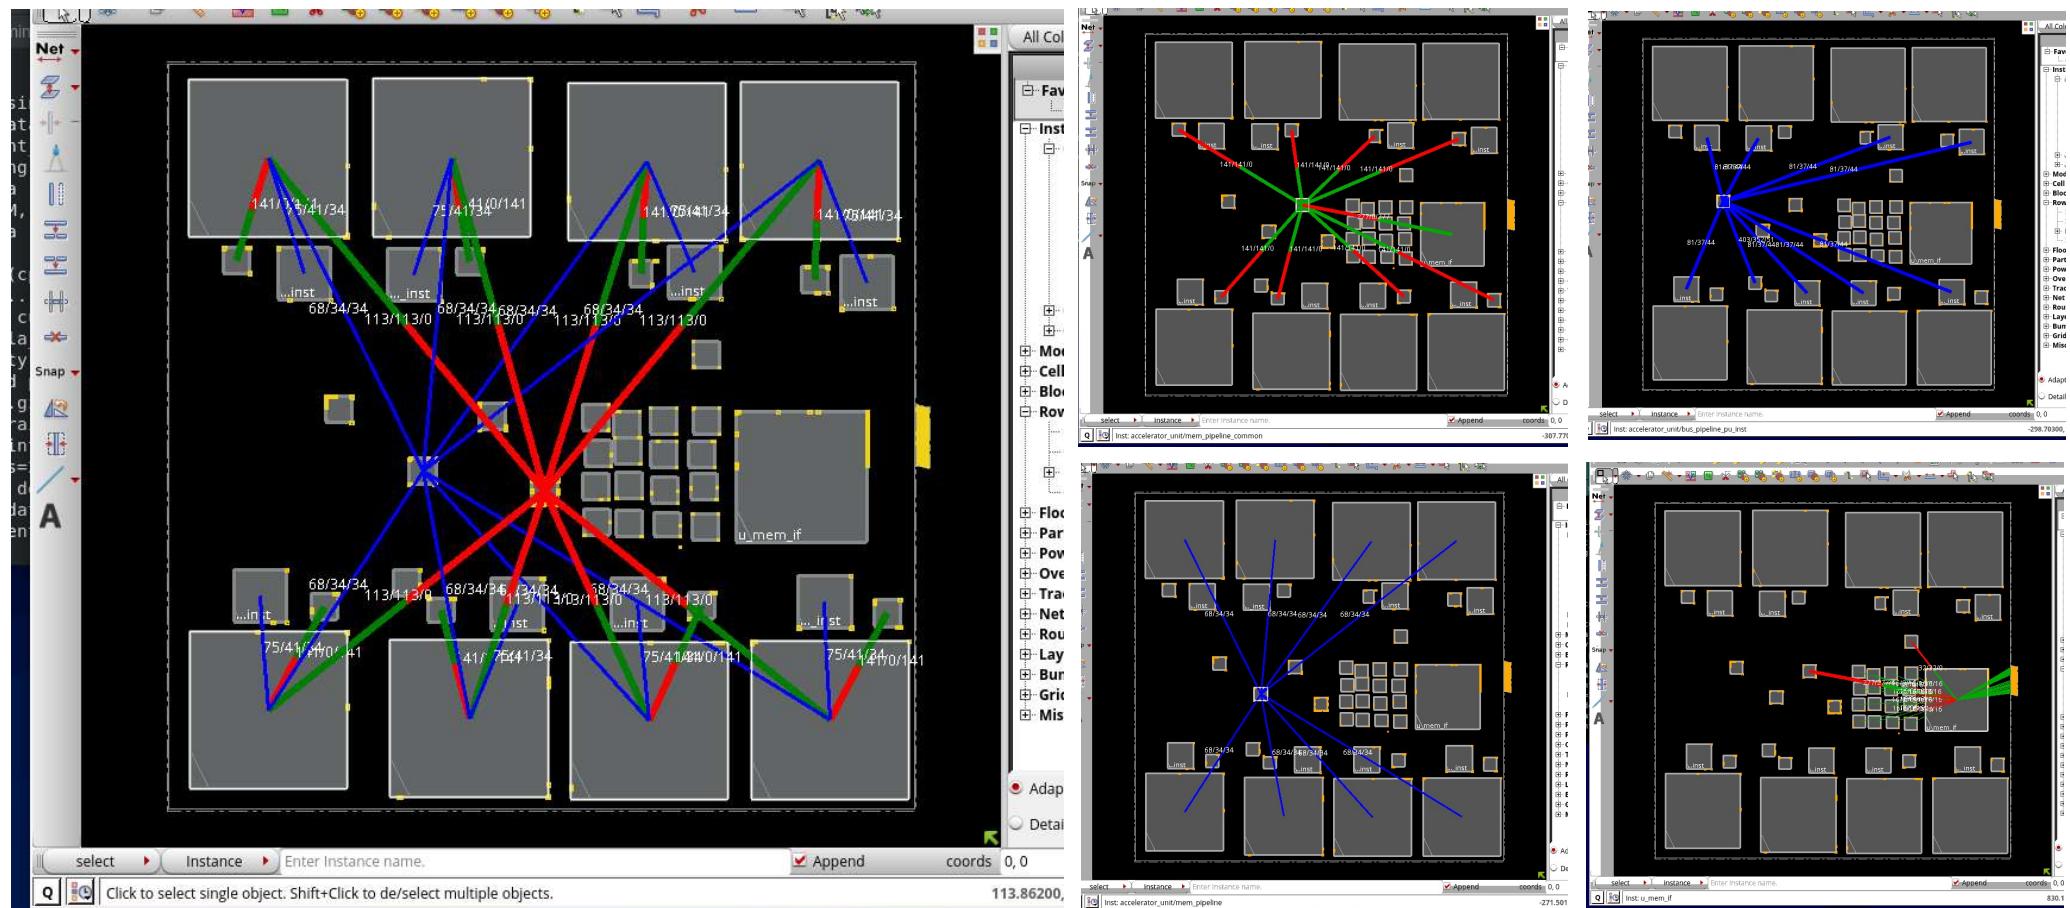


GeneSys: machine learning accelerators for DNN algorithms



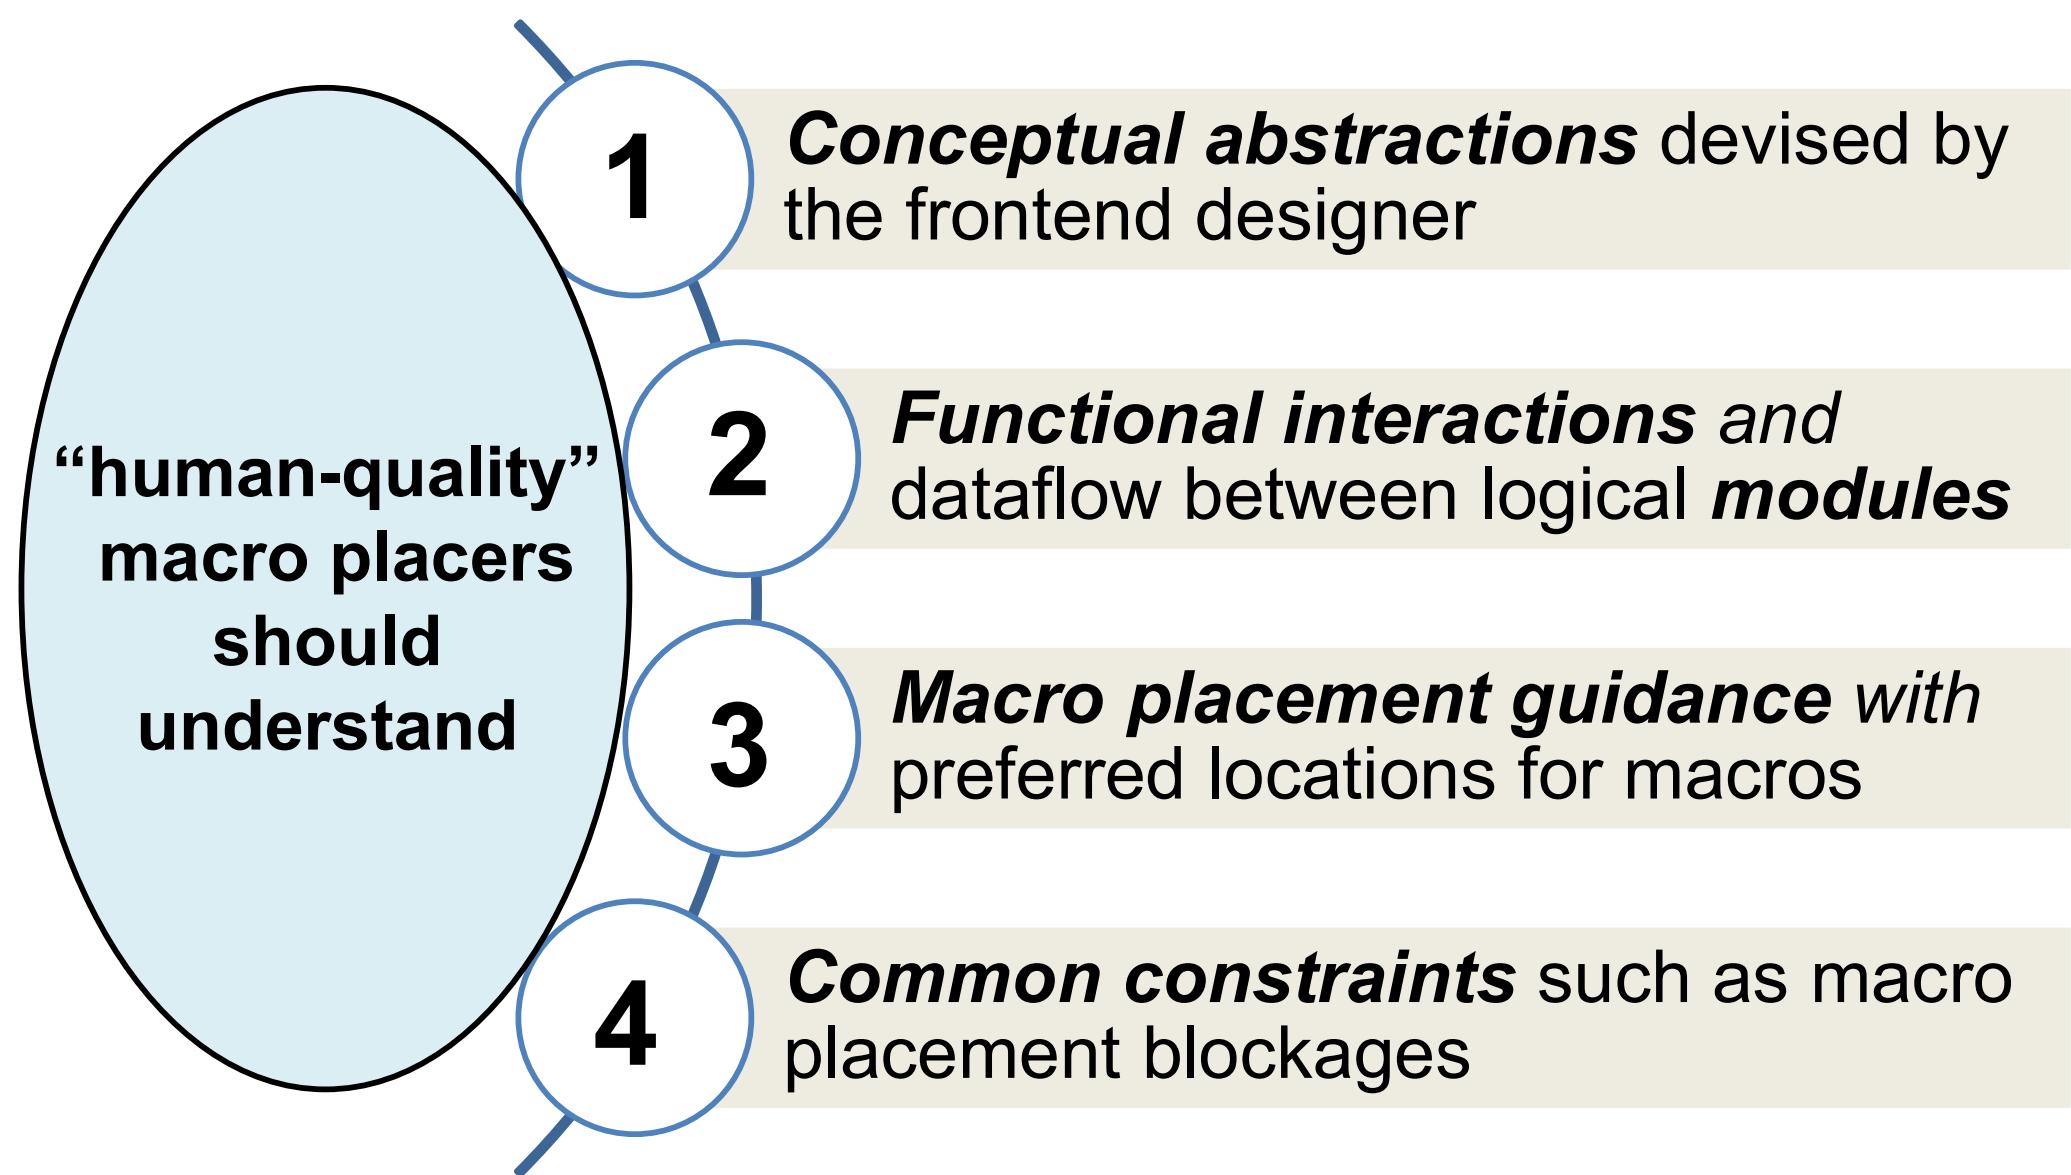
Dataflow Analysis by Humans

- **Dataflow** refers to the movement of data between different functional units.



**TABA: systolic array structures
(Processing Units + Ctrl Logic + Mem Blocks)**

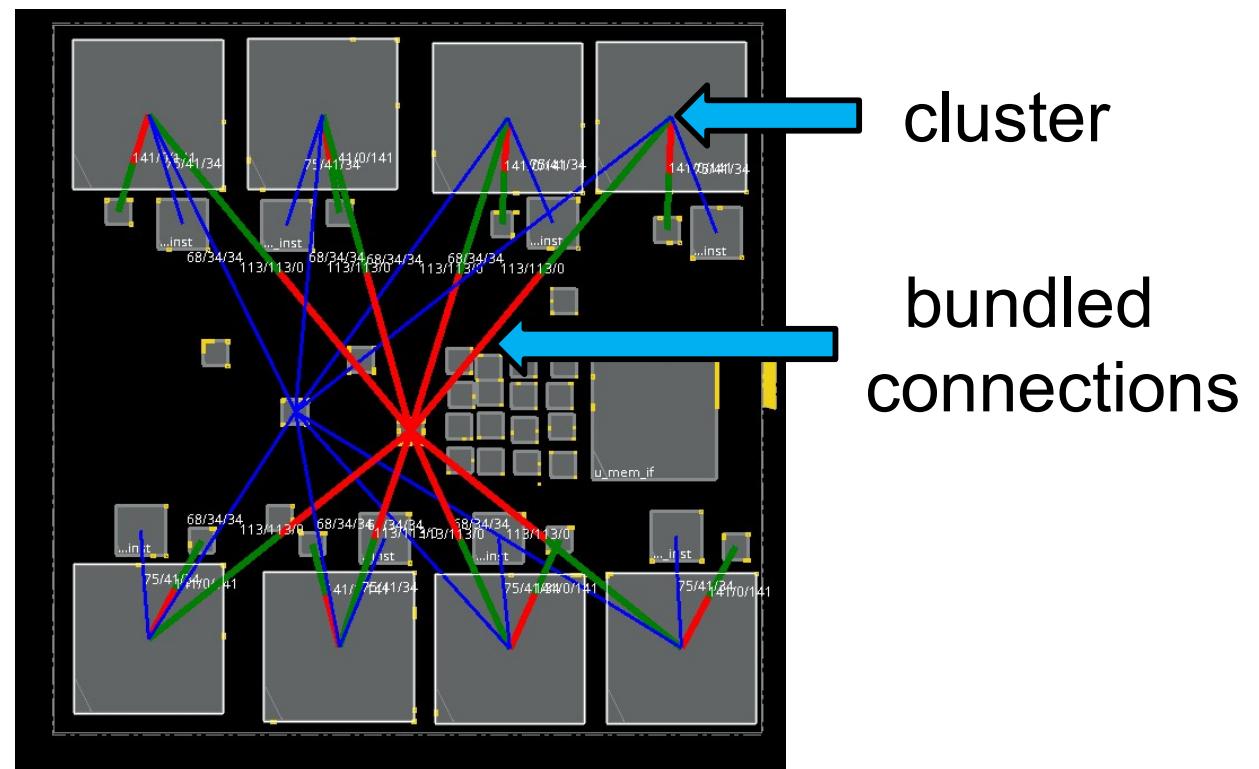
What Should a Human-quality Macro Placer Do ?



Macro Placer in OpenROAD: Hier-RTLMP

- *Conceptual abstractions devised by the frontend designer*
 - Through converting the structural netlist representation of the RTL design into a clustered netlist

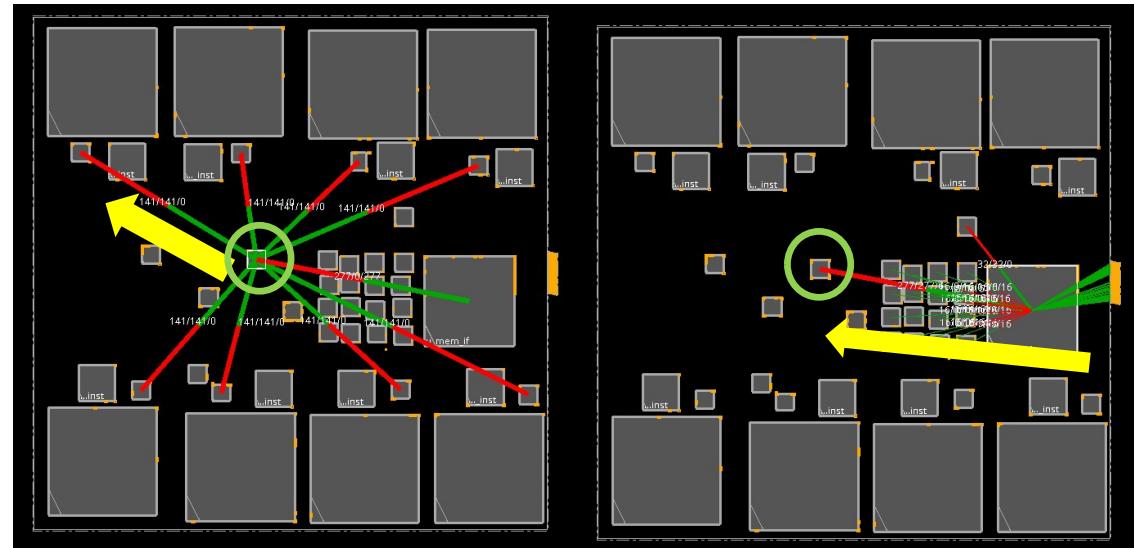
Hier-RTLMP



Macro Placer in OpenROAD: Hier-RTLMP

- *Functional interactions*
 - Through analyzing the dataflow between logical modules (including Input-Output pins)

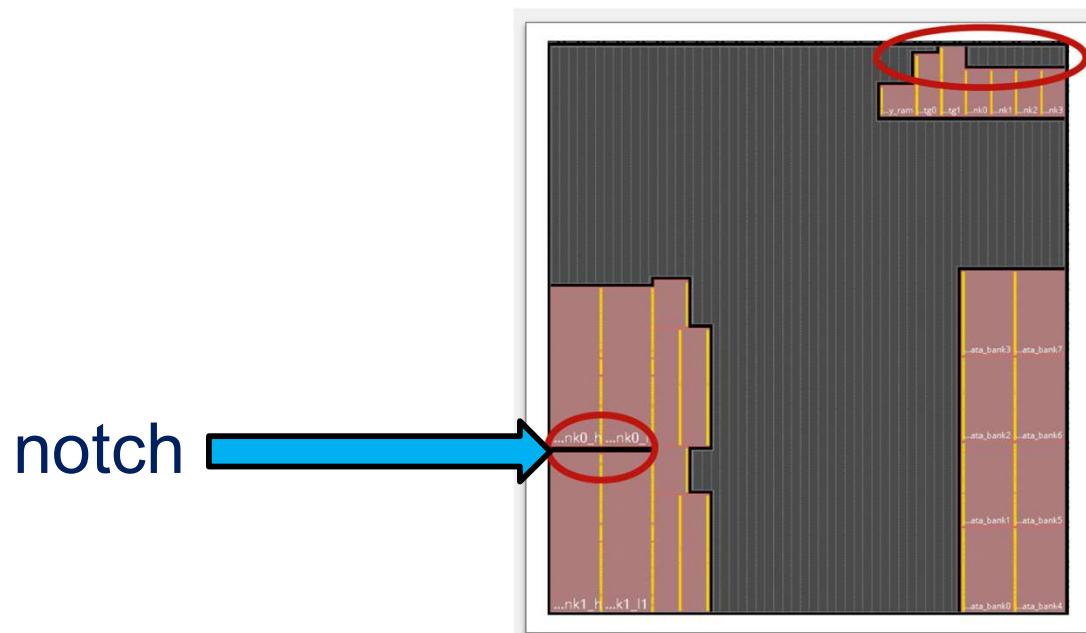
Hier-RTLMP



Macro Placer in OpenROAD: Hier-RTLMP

Hier-RTLMP

- *Macro placement guidance*
 - Through pushing macros to their preferred locations
 - Macro placement blockages
 - Small dead space (notch) avoidance
 - Pushing macros to peripheries

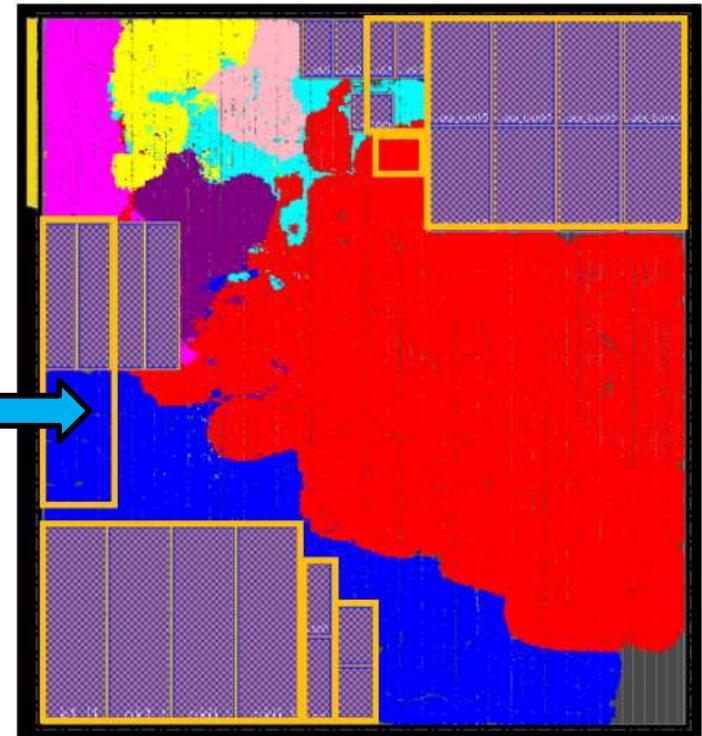


Macro Placer in OpenROAD: Hier-RTLMP

- *Common constraints*
 - Fixed macro constraints
 - Through pushing macros to their preferred locations

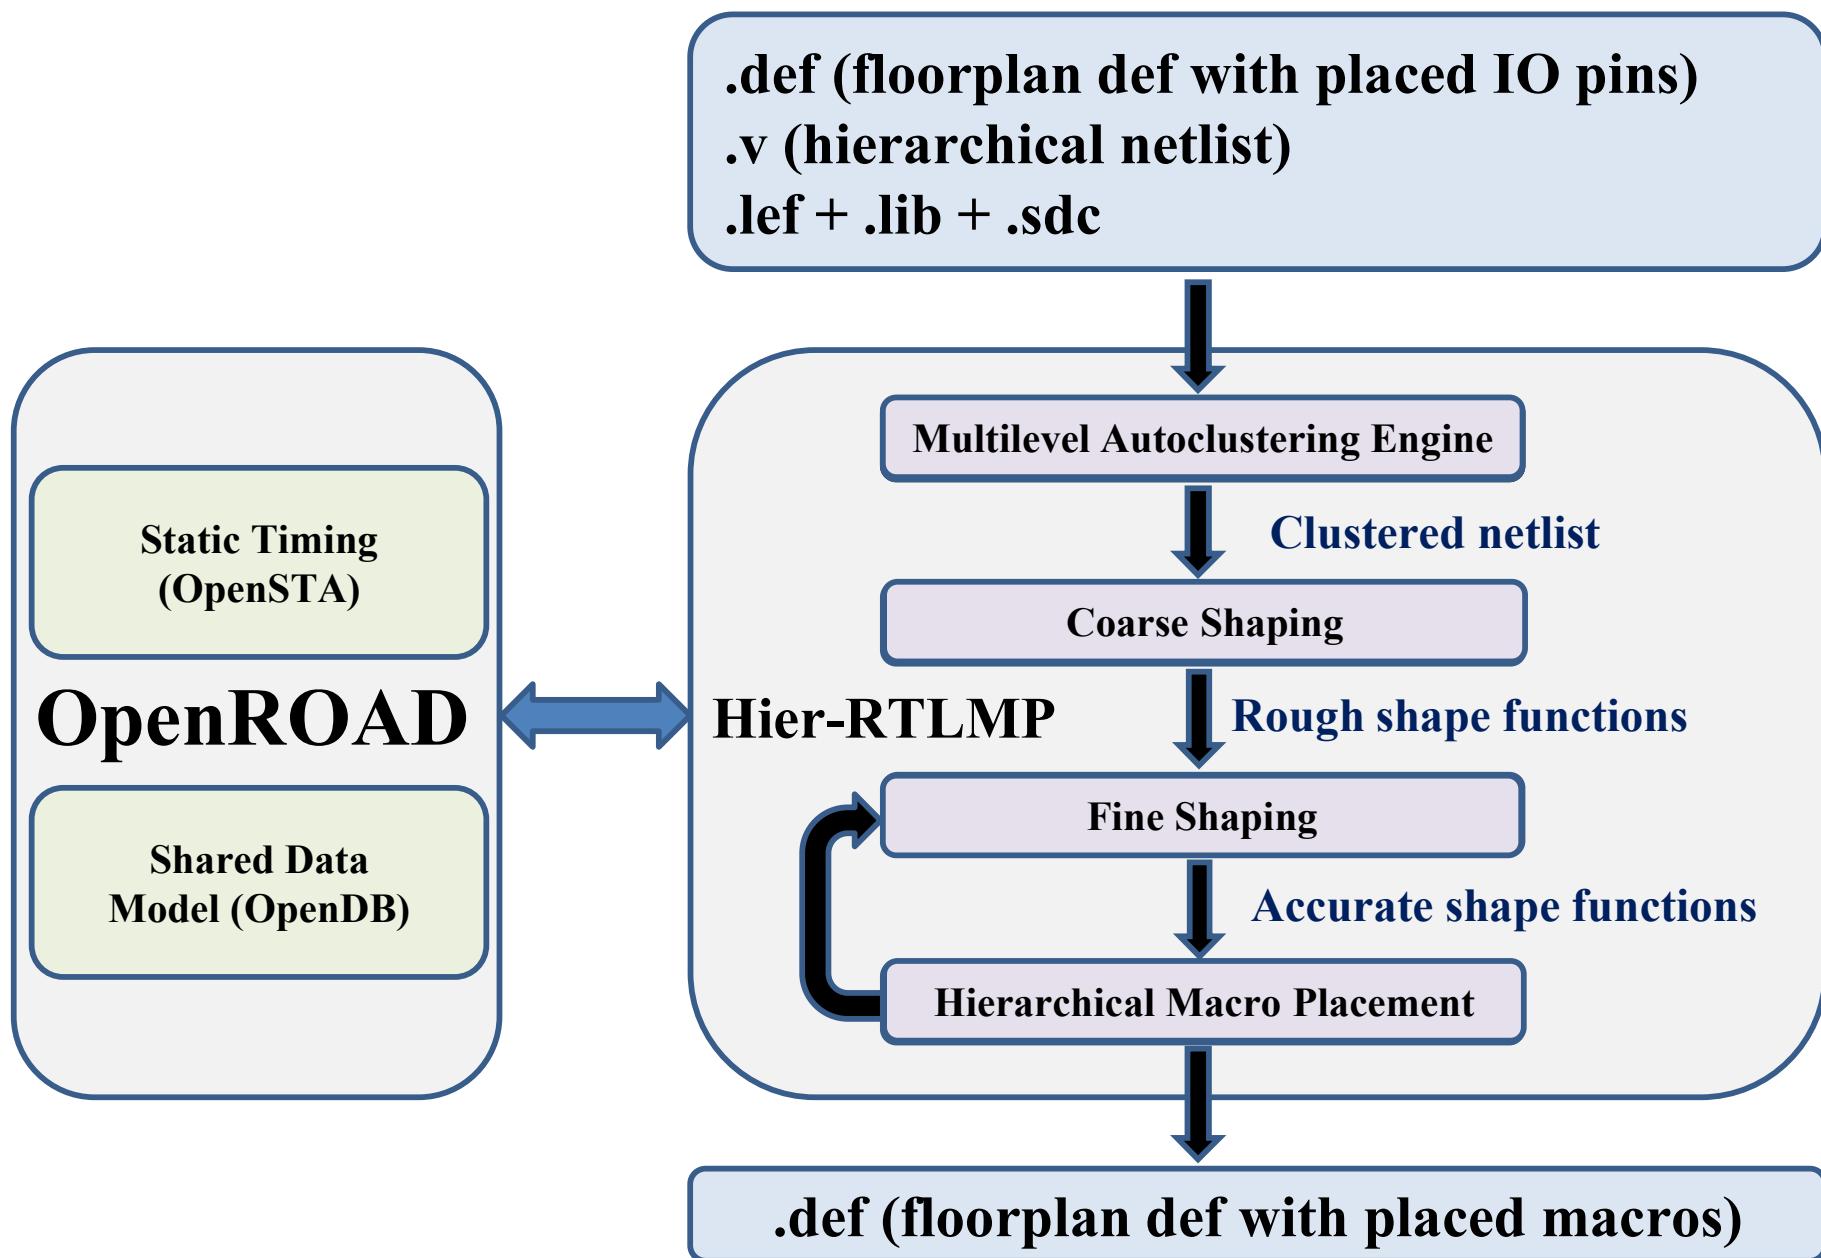
Hier-RTLMP

macro
guidance



Hier-RTLMP flow

(paper: [link](#))



Autoclustering Engine

- Converts the structural netlist representation of the RTL design into a clustered netlist ([clusterEngine.h](#))
 - Logical hierarchy ([code](#))
 - Connection topology between clusters ([code](#))
 - Macro regularity = grouping macros based on footprints ([code](#))
 - Separation between macros and corresponding standard cell logic ([code](#))
 - Dataflow between logical modules (including Input-output pins) ([code](#))

```
283     // Transform the logical hierarchy into a physical hierarchy.
284     void HierRTLMP::runMultilevelAutoclustering()
285     {
286         clustering_engine_ = std::make_unique<ClusteringEngine>(
287             block_, network_, logger_, tritonpart_);
288
289         // Set target structure
290         clustering_engine_->setTree(tree_.get());
291         clustering_engine_->run();
292
293         if (!tree_->has_unfixed_macros) {
294             skip_macro_placement_ = true;
295             return;
296         }
297
298         if (graphics_) {
299             graphics_->finishedClustering(tree_.get());
300         }
301     }
```

Line 284 in hier_rtlmp.cpp ([link](#))



Kahng ECE 260C SP25

```
57     void ClusteringEngine::run()
58     {
59         design_metrics_ = computeModuleMetrics(block_->getTopModule());
60         init();
61
62         if (!tree_->has_unfixed_macros) {
63             return;
64         }
65
66         createRoot();
67         setBaseThresholds();
68
69         mapIOPinsAndPads();
70         createDataFlow();
71
72         createIOClusters();
73         classifyBoundariesStateForIOs();
74
75         if (design_metrics_->getNumStdCell() == 0) {
76             logger_->warn(MPL, 25, "Design has no standard cells!");
77             tree_->has_std_cells = false;
78             treatEachMacroAsSingleCluster();
79         } else {
80             multilevelAutocluster(tree_->root.get());
```

Line 57 in hier_rtlmp.cpp ([link](#))

Coarse Shaping

Wait, what? → SA = how we search over solution representations;
Sequence Pair = how we represent (macro placement) solutions.

- Determine possible rough shapes (area and aspect ratio) for clusters
 - Simulated Annealing + Sequence Pair
 - Cost function: minimize the area

```
329 ✓ void HierRTLMP::runCoarseShaping()
330 {
331     setRootShapes();
332
333     if (tree_->has_only_macros) {
334         logger_->warn(MPL, 27, "Design has only macros!");
335         tree_->root->setClusterType(HardMacroCluster);
336         return;
337     }
338
339     if (graphics_) {
340         graphics_->startCoarse();
341     }
342
343     calculateChildrenTilings(tree_->root.get());
344
345     setPinAccessBlockages();
346     setPlacementBlockages();
347 }
348 }
```

Line 329 in hier_rtlmp.cpp ([link](#))

```
469     // we vary the outline of parent cluster to generate different tilings
470     // we first vary the outline width while keeping outline height fixed
471     // Then we vary the outline height while keeping outline width fixed
472     // Vary the outline width
473     std::vector<float> vary_factor_list{1.0};
474     float vary_step = 1.0 / num_runs_; // change the outline by based on num_runs
475     for (int i = 1; i < num_runs_; i++) {
476         vary_factor_list.push_back(1.0 - i * vary_step);
477     }
478     int remaining_runs = num_runs_;
479     int run_id = 0;
480
481     while (remaining_runs > 0) {
482         SoftSAVector sa_batch;
483         const int run_thread
484             = graphics_ ? 1 : std::min(remaining_runs, num_threads_);
485         for (int i = 0; i < run_thread; i++) {
486             const Rect new_outline(0,
487                                   0,
488                                   outline.getWidth() * vary_factor_list[run_id++],
489                                   outline.getHeight());
490             if (graphics_) {
491                 graphics_->setOutline(micronsToDbu(new_outline));
492             }
493         }
494     }
495 }
```

Simulated Annealing
Line 469 in hier_rtlmp.cpp ([link](#))
[\(SACoreSoftMacro.h\)](#)

```
519     // multi threads
520     std::vector<std::thread> threads;
521     threads.reserve(sa_batch.size());
522     for (auto& sa : sa_batch) {
523         threads.emplace_back(runSA<SACoreSoftMacro>, sa.get());
524     }
525     for (auto& th : threads) {
526         th.join();
527     }
528 }
```

Multi-threading implementation
Line 519 in hier_rtlmp.cpp ([link](#))

Fine Shaping + Hierarchical Macro Placement

- Fine Shaping: determine the accurate shape and aspect ratio for each cluster
 - Pick legal shapes (must fit into the outline of current cluster) for each child clusters

```
1734     const float target_util = target_util_list[run_id];
1735     const float target_dead_space = target_dead_space_list[run_id++];
1736     debugPrint(logger_,
1737             MPL,
1738             "fine_shaping",
1739             1,
1740             "Starting adjusting shapes for children of {}.\n    target_util = "
1741             "{}\n    target_dead_space = {}",
1742             parent->getName(),
1743             target_util,
1744             target_dead_space);
1745     if (!runFineShaping(parent,
1746                         shaped_macros,
1747                         soft_macro_id_map,
1748                         target_util,
1749                         target_dead_space)) {
1750         debugPrint(logger_,
1751                 MPL,
1752                 "fine_shaping",
1753                 1,
1754                 "Cannot generate feasible shapes for children of {}, sa_id: "
1755                 "{}\n    target_util: {}\n    target_dead_space: {}",
1756                 parent->getName(),
1757                 run_id,
1758                 target_util,
1759                 target_dead_space);
1760         continue;
1761     }
```

Line 1731 in hier_rtlmp.cpp ([link](#))

```
1987     // Determine the shape of each cluster based on target utilization
1988     // and target dead space. In contrast to all previous works, we
1989     // use two parameters: target utilization, target_dead_space.
1990     // This is the trick part. During our experiments, we found that keeping
1991     // the same utilization of standard-cell clusters and mixed cluster will make
1992     // SA very difficult to find a feasible solution. With different utilization,
1993     // SA can more easily find the solution. In our method, the target_utilization
1994     // is used to determine the bloating ratio for mixed cluster, the
1995     // target_dead_space is used to determine the bloating ratio for standard-cell
1996     // cluster. The target utilization is based on tiling results we calculated
1997     // before. The tiling results (which only consider the contribution of hard
1998     // macros) will give us very close starting point.
1999     bool HierRTLMP::runFineShaping(Cluster* parent,
2000                                     std::vector<SoftMacro>& macros,
2001                                     std::map<std::string, int>& soft_macro_id_map,
2002                                     float target_util,
2003                                     float target_dead_space)
2004     {
2005         const float outline_width = parent->getWidth();
2006         const float outline_height = parent->getHeight();
2007         float pin_access_area = 0.0;
2008         float std_cell_cluster_area = 0.0;
2009         float std_cell_mixed_cluster_area = 0.0;
2010         float macro_cluster_area = 0.0;
2011         float macro_mixed_cluster_area = 0.0;
2012         // add the macro area for blockages, pin access and so on
2013         for (auto& macro : macros) {
2014             if (macro.getCluster() == nullptr) {
2015                 pin_access_area += macro.getArea(); // get the physical-only area
2016             }
2017         }
```

Line 1999 in hier_rtlmp.cpp ([link](#))

Fine Shaping + Hierarchical Macro Placement

- Hierarchical Macro Placement

- Clusters are placement level by level in a pre-order DFS manner
- Simulated Annealing + Sequence Pair
- Cost function: minimize area, wirelength, penalty functions for different constraints ([code](#))

```
303 ✓ void HierRTLMP::runHierarchicalMacroPlacement()
304 {
305     if (graphics_) {
306         graphics_->startFine();
307     }
308
309     adjustMacroBlockageWeight();
310     placeChildren(tree_->root.get());
311 }
```



```
1101 ✓ void HierRTLMP::placeChildren(Cluster* parent)
1102 {
1103     if (parent->getClusterType() == HardMacroCluster) {
1104         placeMacros(parent);
1105         return;
1106     }
1107
1108     if (parent->isLeaf()) { // Cover IO Clusters && Leaf Std Cells
1109         return;
1110     }
1111
1112     debugPrint(logger_,
1113                 MPL,
1114                 "hierarchical_macro_placement",
1115                 1,
1116                 "Placing children of cluster {}",
1117                 parent->getName());
1118
1119     if (graphics_) {
1120         graphics_->setCurrentCluster(parent);
1121     }
1122 }
```

```
1379     // Note that all the probabilities are normalized to the summation of 1.0.
1380     // Note that the weight are not necessarily summarized to 1.0, i.e., not
1381     // normalized.
1382     std::unique_ptr<SACoreSoftMacro> sa
1383         = std::make_unique<SACoreSoftMacro>(tree_.get(),
1384                                             outline,
1385                                             shaped_macros,
1386                                             placement_core_weights_,
1387                                             boundary_weight_,
1388                                             macro_blockage_weight_,
1389                                             notch_weight_,
1390                                             notch_h_th_,
1391                                             notch_v_th_,
1392                                             pos_swap_prob_ / action_sum,
1393                                             neg_swap_prob_ / action_sum,
1394                                             double_swap_prob_ / action_sum,
1395                                             exchange_swap_prob_ / action_sum,
1396                                             resize_prob_ / action_sum,
1397                                             init_prob_,
1398                                             max_num_step_,
1399                                             num_perturb_per_step,
1400                                             random_seed_,
1401                                             graphics_.get(),
1402                                             logger_);
1403     sa->setNumberOfMacrosToPlace(num_of_macros_to_place);
1404     sa->setCentralizationAttemptOn(true);
1405     sa->setFences(fences);
1406     sa->setGuides(guides);
1407     sa->setNets(nets);
1408     sa->addBlockages(placement_blockages);
1409     sa->addBlockages(macro_blockages);
1410     sa_batch.push_back(std::move(sa));
1411 }
```

Line 303 in hier_rtlmp.cpp ([link](#))



Line 1379 in hier_rtlmp.cpp ([link](#))

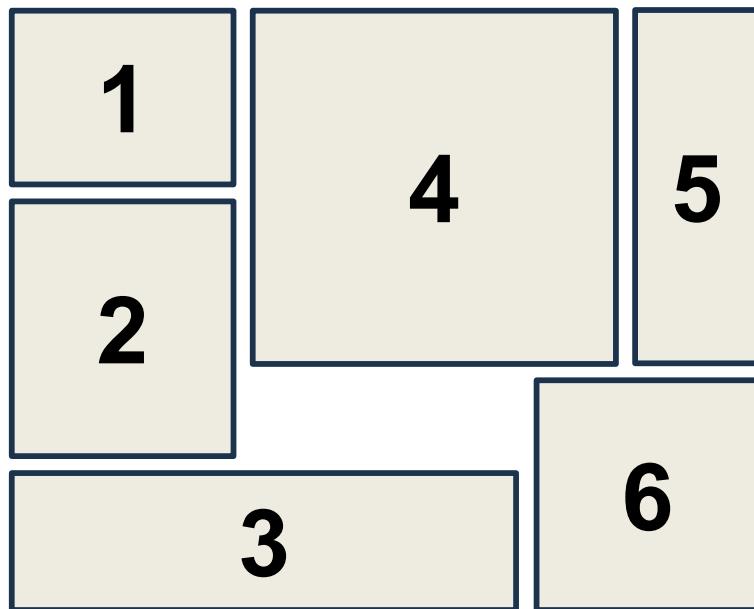
Sequence Pair and Simulated Annealing

Sequence Pair: How we represent the floorplan solution

Simulated Annealing: How we (heuristically) optimize the floorplan

Sequence Pair (SP) (paper: [link](#))

- SP consists of an ordered pair (Γ^+, Γ^-) of cluster sequences.
- Given a sequence pair (Γ^+, Γ^-) , the location of clusters can be determined as follows:
 - (**Horizontal constraint**): cluster i is left to cluster j if i appears before j in both Γ^+ and $\Gamma^- \rightarrow (\dots i \dots j \dots, \dots j \dots i \dots)$
 - (**Vertical constraint**): cluster i is below cluster j if i appears after j in Γ^+ and i appears before j in $\Gamma^- \rightarrow (\dots j \dots i \dots, \dots j \dots i \dots)$



Longest Common Subsequence !
(paper: [link](#))

$$(\Gamma^+, \Gamma^-) = (124536, 326145)$$

From Floorplan to Sequence Pair

- A **sequence pair** consists of two module name sequences. For example, the floorplan shown at right can be represented as the sequence pair $(\Gamma_+, \Gamma_-) = (124536, 326145)$.

- **Construction of Γ_+**

- For each module, draw a **right-up locus** and a **left-down locus** as shown in Figure (a). Then we order these loci from the left to right. Γ_+ is resulting order of module names.

- **Construction of Γ_-**

- For each module, we draw a **up-left locus** and a **down-right locus** as shown in figure (b). Then we order these loci from the left to right. Γ_- is the resulting order of module names.

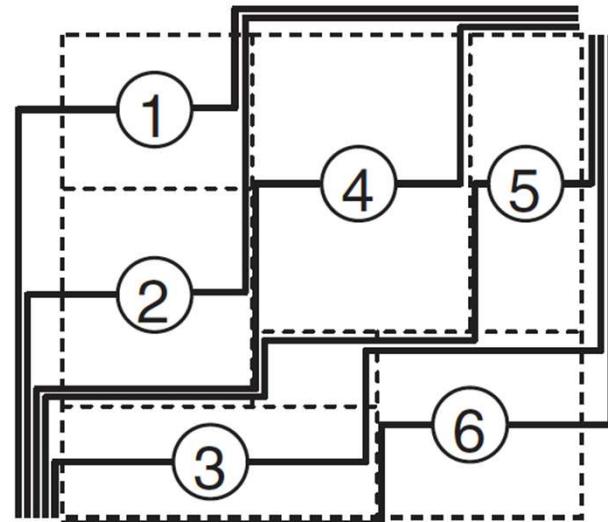
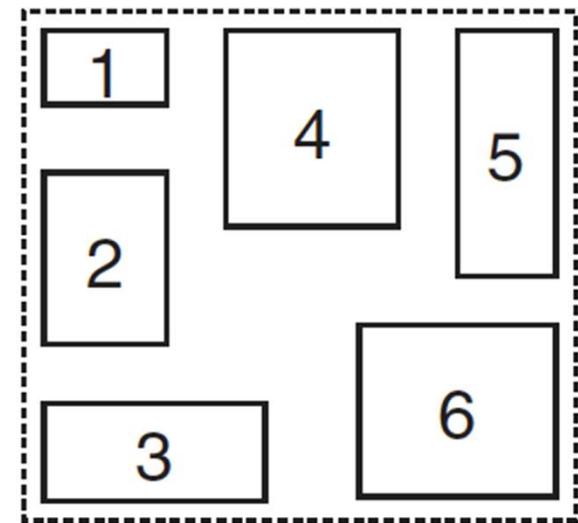


Figure (a)

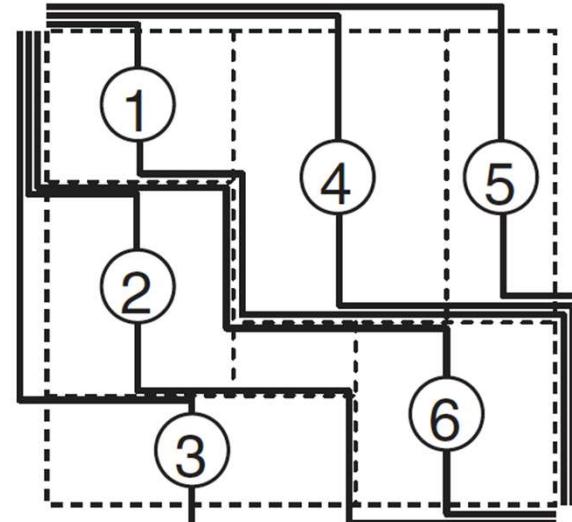


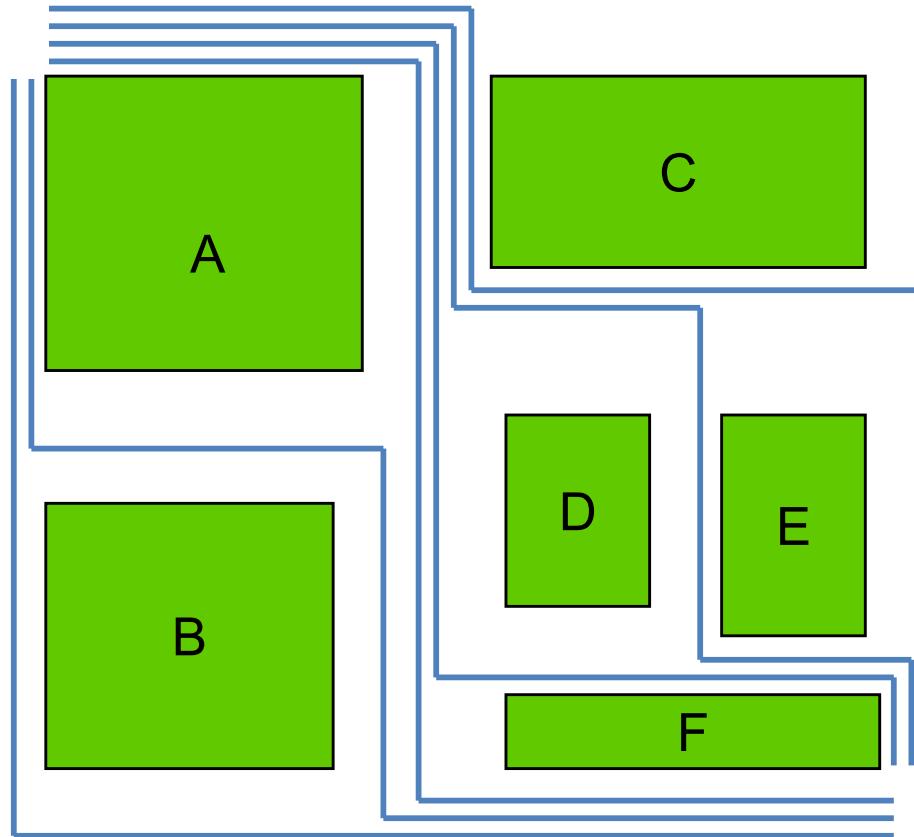
Figure (b)

Think: “non-crossing staircases”

Another example of (up-left, down-right) ordering Γ_-

$\Gamma_- = \text{CEDFAB}$

(Exercise: What is Γ_+ ?)



From Sequence Pair to Floorplan

- Given an SP (Γ_+, Γ_-) , the geometric relation between modules (clusters) is derived from the SP:
 - Horizontal constraint:** module i is to the left of module j if i appears **before** j in both Γ_+ and Γ_- ($\dots i \dots j \dots, \dots i \dots j \dots$)
 - Vertical constraint:** module i is below module j if i appears **after** j in Γ_+ and i appears **before** j in Γ_- ($\dots j \dots i \dots, \dots i \dots j \dots$)
- From horizontal constraints: Create a **horizontal constraint graph** with a source and a sink, and a node-weighted directed acyclic graph $G_H(V, E)$, where V is the set of nodes, and E is the set of edges as follows:
 - V source s, sink t, and n nodes labeled with module names
 - E: (s, t) and (i, t) for each module i, and (i, j) if and only if module i is on the left of module j (horizontal constraint).
 - Node weight : zero for s and t, width of module i for node i.
- From vertical constraints: create a vertical constraint graph $G_V(V, E)$.

From Sequence Pair to Floorplan

- For sequence pair $(\Gamma_+, \Gamma_-) = (124536, 326145)$, we construct the HCG (a) and the VCG (b)
- A “longest common subsequences” algorithm (see AdyaM03) is used in calculating this final floorplan (c)

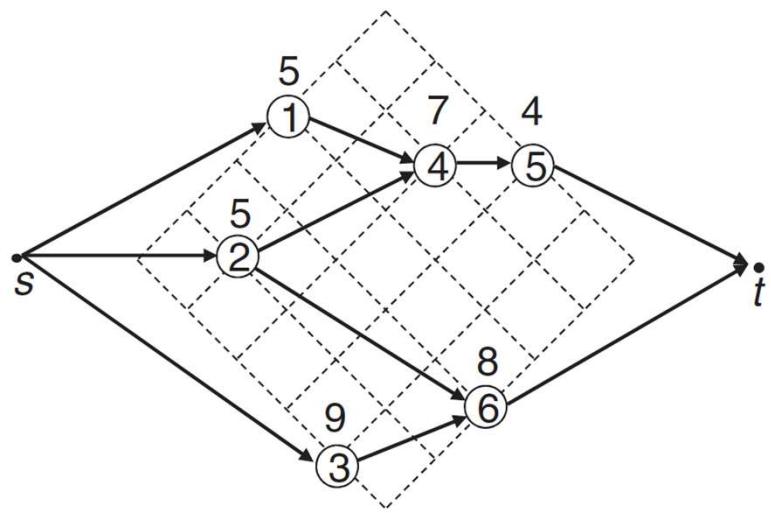


Figure (a)

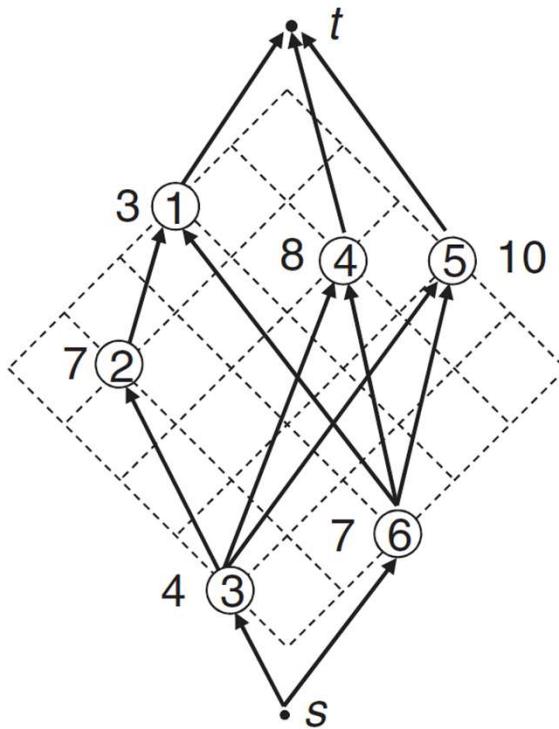


Figure (b)

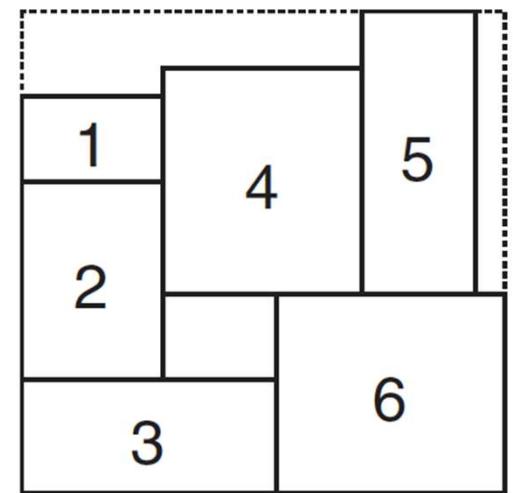


Figure (c)

Sequence Pair (SP) (code)

```
479 // Determine the positions of macros based on sequence pair
480 template <class T>
481 void SimulatedAnnealingCore<T>::packFloorplan()
482 {
483     for (auto& macro_id : pos_seq_) {
484         macros_[macro_id].setX(0.0);
485         macros_[macro_id].setY(0.0);
486     }
487
488     // Each index corresponds to a macro id whose pair is:
489     // <Position in Positive Sequence , Position in Negative Sequence>
490     std::vector<std::pair<int, int>> sequence_pair_pos(pos_seq_.size());
491
492     // calculate X position
493     for (int i = 0; i < pos_seq_.size(); i++) {
494         sequence_pair_pos[pos_seq_[i]].first = i;
495         sequence_pair_pos[neg_seq_[i]].second = i;
496     }
497
498     std::vector<float> accumulated_length(pos_seq_.size(), 0.0);
499     for (int i = 0; i < pos_seq_.size(); i++) {
500         const int macro_id = pos_seq_[i];
501         const int neg_seq_pos = sequence_pair_pos[macro_id].second;
502
503         macros_[macro_id].setX(accumulated_length[neg_seq_pos]);
504
505         const float current_length
506             = macros_[macro_id].getX() + macros_[macro_id].getWidth();
507
508         for (int j = neg_seq_pos; j < neg_seq_.size(); j++) {
509             if (current_length > accumulated_length[j]) {
510                 accumulated_length[j] = current_length;
511             } else {
512                 break;
513             }
514         }
515     }
516
517     width_ = accumulated_length[pos_seq_.size() - 1];
```

```
519     // calculate Y position
520     std::vector<int> reversed_pos_seq(pos_seq_.size());
521     for (int i = 0; i < reversed_pos_seq.size(); i++) {
522         reversed_pos_seq[i] = pos_seq_[reversed_pos_seq.size() - 1 - i];
523     }
524
525     for (int i = 0; i < pos_seq_.size(); i++) {
526         sequence_pair_pos[reversed_pos_seq[i]].first = i;
527         sequence_pair_pos[neg_seq_[i]].second = i;
528
529         // This is actually the accumulated height, but we use the same vector
530         // to avoid more allocation.
531         accumulated_length[i] = 0.0;
532     }
533
534     for (int i = 0; i < pos_seq_.size(); i++) {
535         const int macro_id = reversed_pos_seq[i];
536         const int neg_seq_pos = sequence_pair_pos[macro_id].second;
537
538         macros_[macro_id].setY(accumulated_length[neg_seq_pos]);
539
540         const float current_height
541             = macros_[macro_id].getY() + macros_[macro_id].getHeight();
542
543         for (int j = neg_seq_pos; j < neg_seq_.size(); j++) {
544             if (current_height > accumulated_length[j]) {
545                 accumulated_length[j] = current_height;
546             } else {
547                 break;
548             }
549         }
550     }
551
552     height_ = accumulated_length[pos_seq_.size() - 1];
553
554     if (graphics_) {
555         graphics_->saStep(macros_);
556     }
557 }
```

Line 479 in SimulatedAnnealingCore.cpp ([link](#))

Constraints in Hierarchical Macro Placement

- **Fixed outline:** All clusters should be placed within the fixed outline specified by users. [\(code\)](#)
- **Macro peripheral bias:** All macros should be pushed to peripheries as much as possible. [\(code\)](#)
- **Macro blockage:** All macros should not overlap with macro placement blockages. [\(code\)](#)
- **Pin access:** All macros should be kept from blocking access of input-output pins. [\(code\)](#)
- **Macro guidance:** All macros should be placed near specified regions if users provide such constraints. [\(code\)](#)
- **Notch avoidance:** A decent floorplan should avoid “dead space” which cannot be used effectively by P&R tools. [\(code\)](#)

Simulated Annealing (SA)

- **Kirkpatrick, Gelatt, Vecchi, Science (1983):** One of the most cited scientific papers ever
- SA is one of many “metaheuristics” that are used to deal with instances of intractable (NP-hard) combinatorial problems
 - Genetic algorithms (Holland, U. Michigan)
 - Tabu search (Glover, U. Colorado)
 - ...
- Combinatorial optimization has a physical analogy to the annealing (slow cooling) of metals to produce a perfectly-ordered, minimum-energy state: a “state” is a “solution”, “energy” is “cost”, etc.
- **Basic idea**
 - **Initialize** – Start with a random initial solution. Initialize high “temperature”.
 - **Step 2: “Move Gen” – Perturb** current solution to obtain a ‘neighbor’ solution
 - **Step 3: Calculate cost change** – calculate the change in solution cost due to the move (minimization: negative change is better, positive change is worse)
 - **Step 4: Accept/Reject** – Depending on the cost change, accept or reject the move. Probability of acceptance depends on current “temperature”.
 - **Step 5: Update** – Update temperature, current solution. **Go to Step 2.**
 - **Continue until termination** condition (‘freezing’ or ‘quenching’) is satisfied

Neighborhood Structure

- Topology (i.e., “graph”) over the space of solutions
- Adjacency in this graph is induced by the definition of “perturbation” (aka “move”)
- Traveling Salesperson Problem: switch positions of two cities in the tour
- Graph Bisection Problem: swap a pair of vertices between the partitions
- TSP: “optimally rearrange the positions of four cities in the tour” → more neighbors, more powerful move, more work to generate a move

SA Pseudocode

<http://www.ecs.umass.edu/ece/labs/vlsicad/ece665/slides/SimulatedAnnealing.ppt>

Algorithm SIMULATED-ANNEALING

Begin

temp = INIT-TEMP;

currentSol = INIT-SOLUTION;

for i = 1 to M

candidateSol = NEIGHBOR(*currentSol*);

ΔC = COST(*candidateSol*) – COST(*currentSol*);

if ($\Delta C < 0$) **then**

currentSol = *candidateSol*;

else with $Pr = e^{-(\Delta C / temp)}$

currentSol = *candidateSol*;

temp = SCHEDULE(*temp*);

End

What happens when $temp = +\infty$?

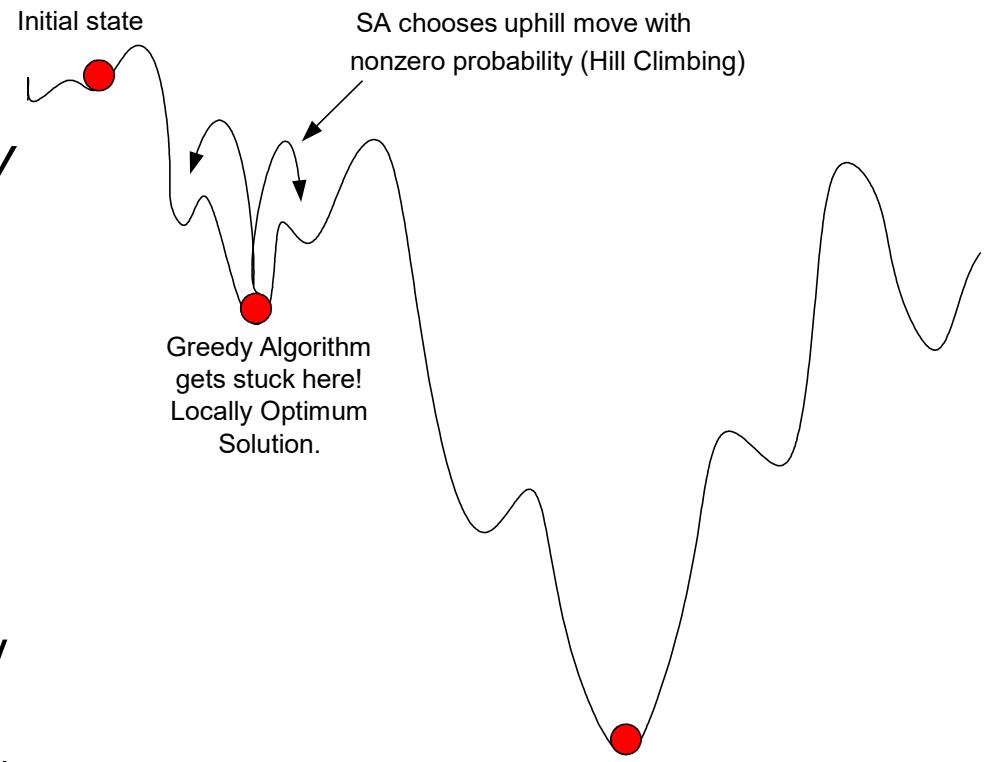
What happens when $temp = 0$?

Simulated Annealing Facts

- NEIGHBOR(solution) defines a **topology** over all solutions in the solution space
- At a **fixed** value of **temp**, SA behavior corresponds to a *homogeneous Markov chain*
 - Matrix of transition probabilities between states
- Steady-state (= equilibrium) probability of the Markov chain being in state A is proportional to $e^{(-\text{cost}(A)/\text{temp})}$
 - When $\text{temp} \rightarrow 0$, exponentially more likely to be in the global optimum state
 - “SA is optimal” (in the limit of ‘infinite time’)
 - Of course, we spend only a finite amount of time (#moves) at any temperature value
 - **Is cooling the best strategy with finite time?** See Boese/Kahng, 1993 [[link](#)] [[link](#)]

Question: Why is SA so widely applied?

Answer: (1) Easy to implement. (2) Can “handle” almost any discrete optimization (search in a solution space) task. [But: (3) “discrete”, “almost any” are two-edged swords. And (4) SA needs help to scale ...]



SA converges to global opt solution with $\text{Pr} = 1$
(in limit of infinite time, infinitely slow cooling)

Neighborhood Structure of Sequence Pair

- Solution space
 - For n modules, the lengths of Γ_+ and Γ_- are both n , and thus each of Γ_+ and Γ_- have $n!$ permutations. There are $(n!)^2$ permutations for a sequence pair with n modules.
- Neighborhood structure
 - M1 Rotate a module
 - M2 Swap two modules in only one sequence
 - M3: Swap two modules in both sequences

Simulated Annealing ([SimulatedAnnealingCore.h](#))

```
759     void SimulatedAnnealingCore<T>::fastSA()
760     {
761         float cost = calNormCost();
762         float pre_cost = cost;
763         float delta_cost = 0.0;
764         int step = 1;
765         float temperature = init_temperature_;
766         const float min_t = 1e-10;
767         const float t_factor
768             = std::exp(std::log(min_t / init_temperature_) / max_num_step_);
769
770         // Used to ensure notch penalty is used only in the latter steps
771         // as it is too expensive
772         notch_weight_ = 0.0;
773
774         int num_restart = 1;
775         const int max_num_restart = 2;
776
777         if (isValid()) {
778             updateBestValidResult();
779         }
780
781         while (step <= max_num_step_) {
782             for (int i = 0; i < num_perturb_per_step_; i++) {
783                 perturb();
784                 cost = calNormCost();
785
786                 const bool keep_result
787                     = cost < pre_cost
788                         || best_valid_result_.sequence_pair.pos_sequence.empty();
789                 if (isValid() && keep_result) {
790                     updateBestValidResult();
791                 }
792             }
793         }
794     }
```

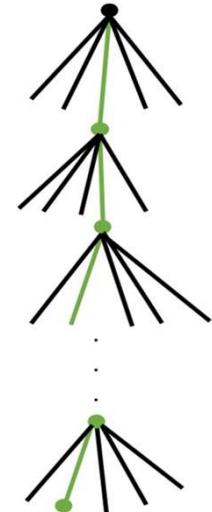
Line 759 in [SimulatedAnnealingCore.cpp](#) ([link](#))

```
793         delta_cost = cost - pre_cost;
794         const float num = distribution_(generator_);
795         const float prob
796             = (delta_cost > 0.0) ? std::exp((-1) * delta_cost / temperature) : 1;
797         if (num < prob) {
798             pre_cost = cost;
799         } else {
800             restore();
801         }
802     }
803
804     temperature *= t_factor; Temperature update
805     step++;
806
807     cost_list_.push_back(pre_cost);
808     T_list_.push_back(temperature);
809
810     if ((num_restart <= max_num_restart)
811         && (step == std::floor(max_num_step_ / max_num_restart))
812             && (outline_penalty_ > 0.0))) {
813         shrink();
814         packFloorplan();
815         calPenalty();
816         pre_cost = calNormCost();
817         num_restart++;
818         step = 1;
819         num_perturb_per_step_ *= 2;
820         temperature = init_temperature_;
821     }
822
823     if (step == max_num_step_ - macros_.size() * 2) {
824         notch_weight_ = original_notch_weight_;
825         packFloorplan();
826         calPenalty();
827         pre_cost = calNormCost();
828     }
829 }
830
831 packFloorplan();
832 if (graphics_) {
833     graphics_->doNotSkip();
834 }
835 calPenalty();
836
837 if (!isValid() && !best_valid_result_.sequence_pair.pos_sequence.empty()) {
838     useBestValidResult();
839 }
840 }
```

Accept or Reject ?

Go-With-The-Winners

- D. Aldous and U. Vazirani, “Go With the Winners Algorithms”, Proc. FOCS, 1994, pp. 492-501. [[link](#)]
 - Launch multiple optimization threads
 - Periodically identify the most promising thread(s)
 - Clone the promising thread(s) and terminate others
 - (*continue with the optimization*) (cf. “evolutionary”, “PPSN”)
- Implementation: [here](#)
 - Initialize all SA workers in parallel
 - **Main Loop:**
 - Run each SA worker for *sync_iter iterations*, in parallel
 - Stop if *Iter* iterations have been performed; otherwise, select top *k* workers and replicate their solutions to remaining workers
 - Write out best macro placement solution of each worker



```
iter_count ← 0;  
sync_iter ← Iters × sync_freq;  
while true do  
    end_iter ← min(Iter, iter_count + sync_iter);  
    for i ← 0 to W - 1 in parallel do  
        | Each worker performs (end_iter - iter_count) SA  
        | iterations; applying N × #macro moves per  
        | iteration and updating temperature;  
        iter_count ← end_iter;  
        if iter_count = Iters then  
            | break;  
        candidate_solutions ← extractTopK(workers, k);  
        Evenly distribute these top-k solutions across all the  
        workers;
```

Write out the best solution of each worker.

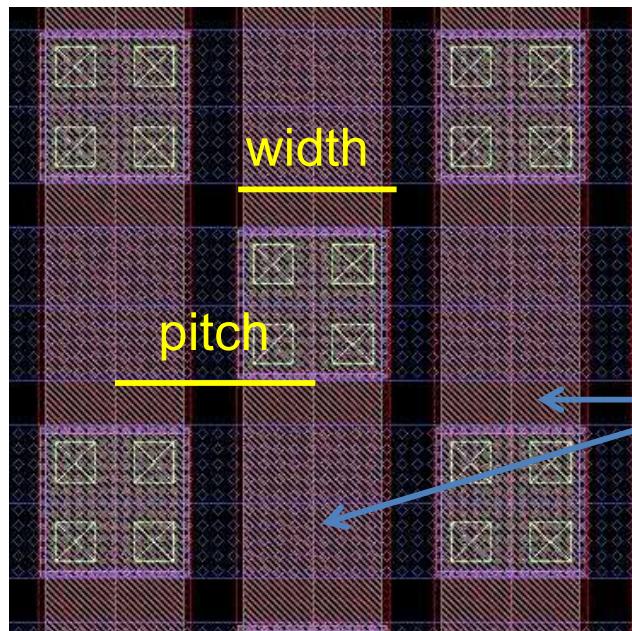
Power Grid Generation

Thanks to Dr. Peter Gadfort ([LinkedIn](#))

Director of Silicon Engineering, Zero ASIC

Power Grid Design and Analysis

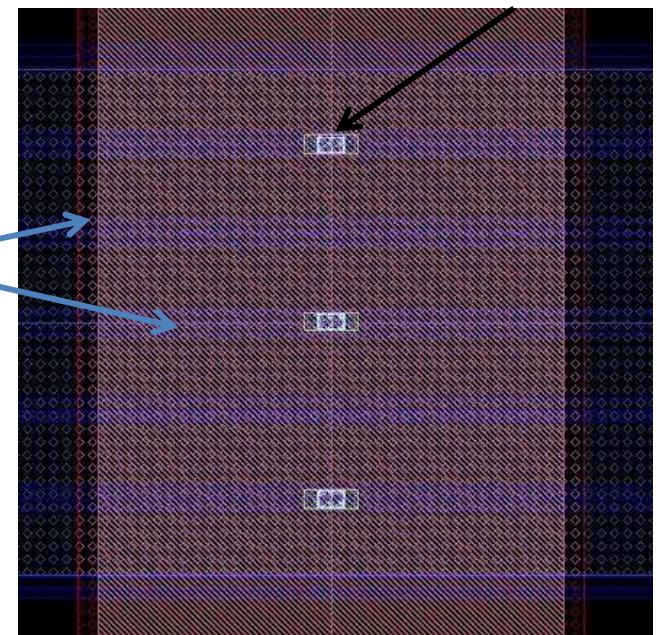
- Typically, power grids are constructed on higher metal layers due to their lower sheet resistance and higher current density specification. Often top metal layers are thick.
- Unless upper metal layers are used for signal routing, the usual approach is to **maximize the power grid**, so long as the maximum metal density rules are satisfied.
- We need to determine the following parameters.
 - Width and pitch of the power mesh stripes (depends on peak current)
 - How frequently to tap down from primary mesh to M1 rails
 - Via stack (e.g. – 1x4, 2x2, or 2x3), should be chosen carefully to avoid blocking extra routing tracks and give low resistance
- The two primary concerns are IR drop and electromigration.



Power Rails

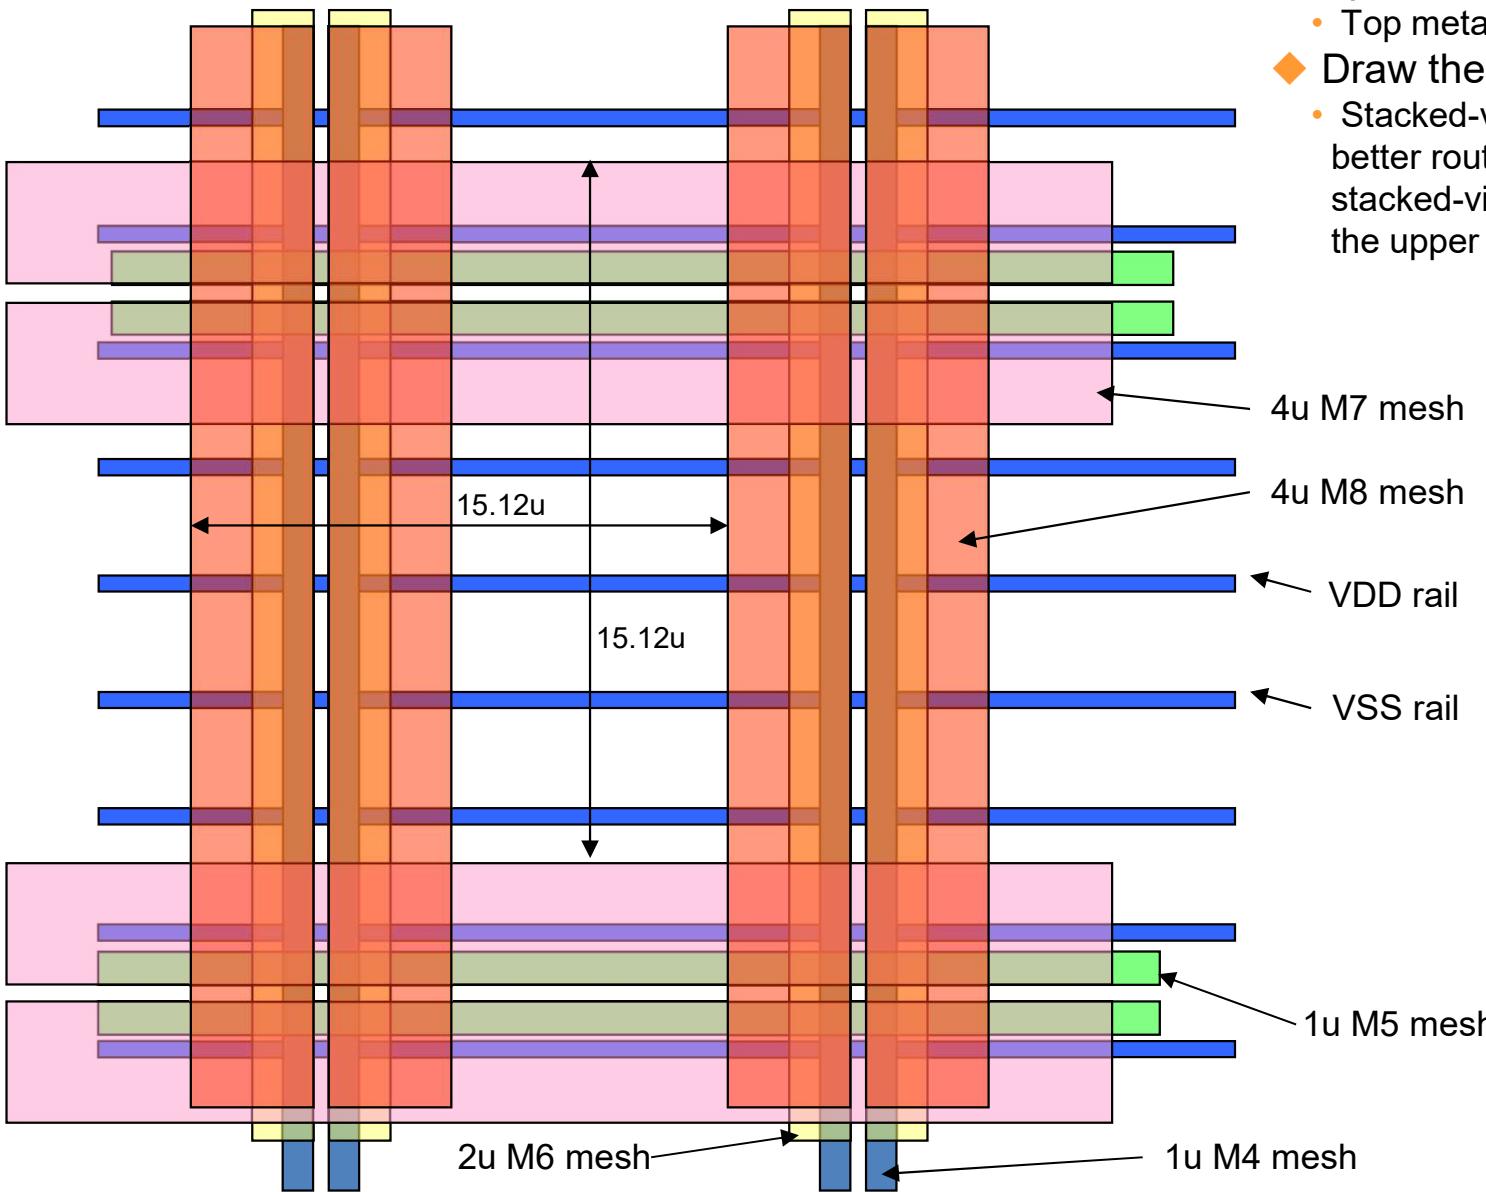
Power Stripes

Via Stack (2x3) 



Power Grid Design and Analysis

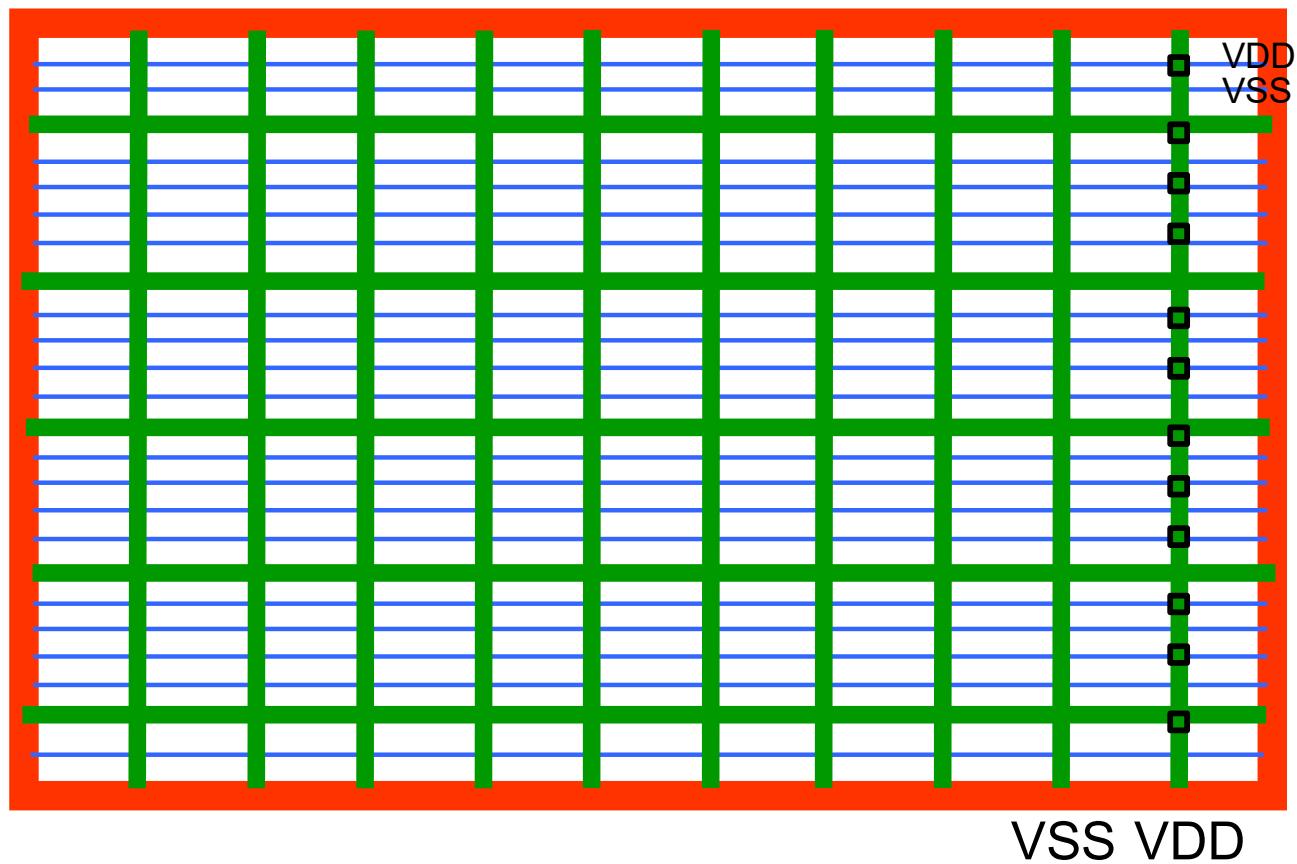
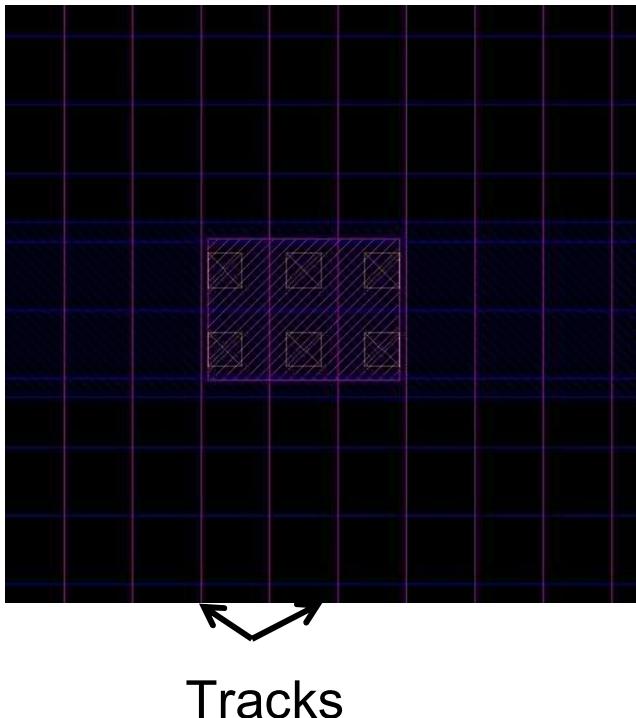
- Structure of the multi-layer power mesh (You can see detail process using slide-show)
 - Space of VDD mesh and VSS mesh
 - Top metal layer's minimum space rule
 - Draw the lowest mesh first
 - Stacked-via range must be changed to get better routability. It is recommended that the stacked-via size is the cross sectional area of the upper and the lower metal of the via.



M1 rail:
M4 mesh:
M5 mesh:
M6 mesh:
M1_2B mesh:
M2_2B mesh:

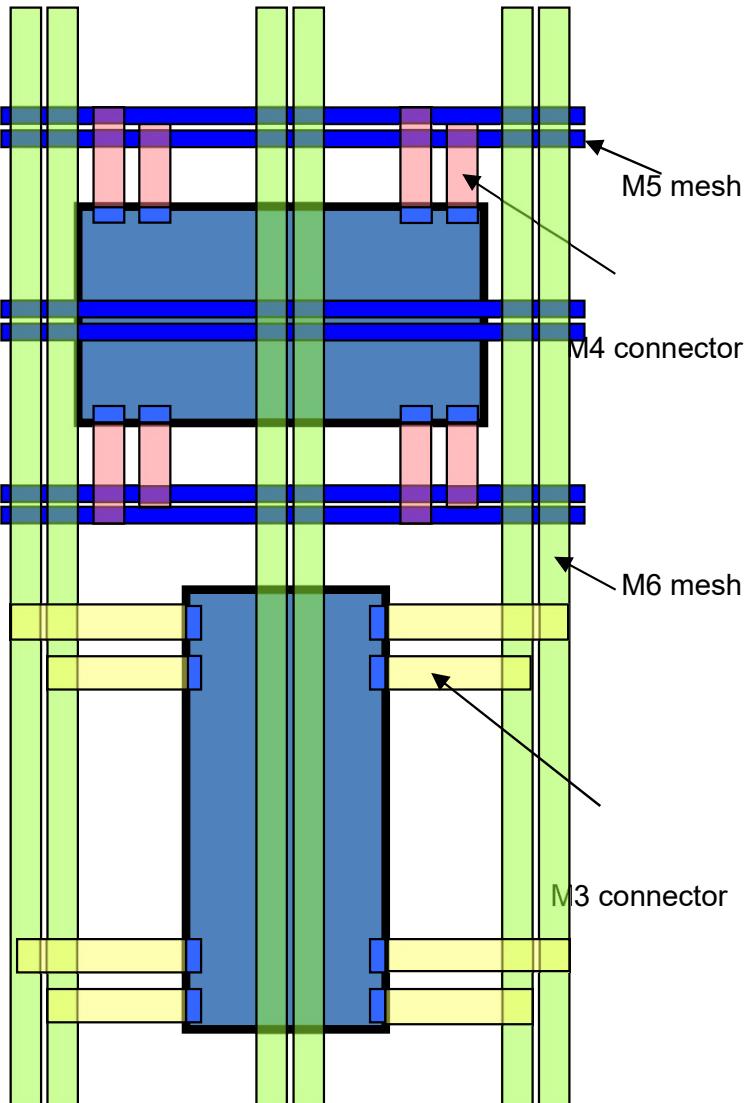
Power Grid Design and Analysis

- Via Stacks
 - Typically we drop vias every time an upper layer power stripe crosses an M1 rail.
 - We choose the via stack so that the M1 rail is the limiting factor for IR drop – not the via stack.
 - The via stack can influence routability, depending on the number of tracks that it blocks.
 - Below example has 2x3 via array and blocks 4 vertical + 3 horizontal tracks.
 - A 2x2 array would block 3 vertical + 3 horizontal tracks.
 - A 1x4 array would block 5 vertical + 1 horizontal tracks.

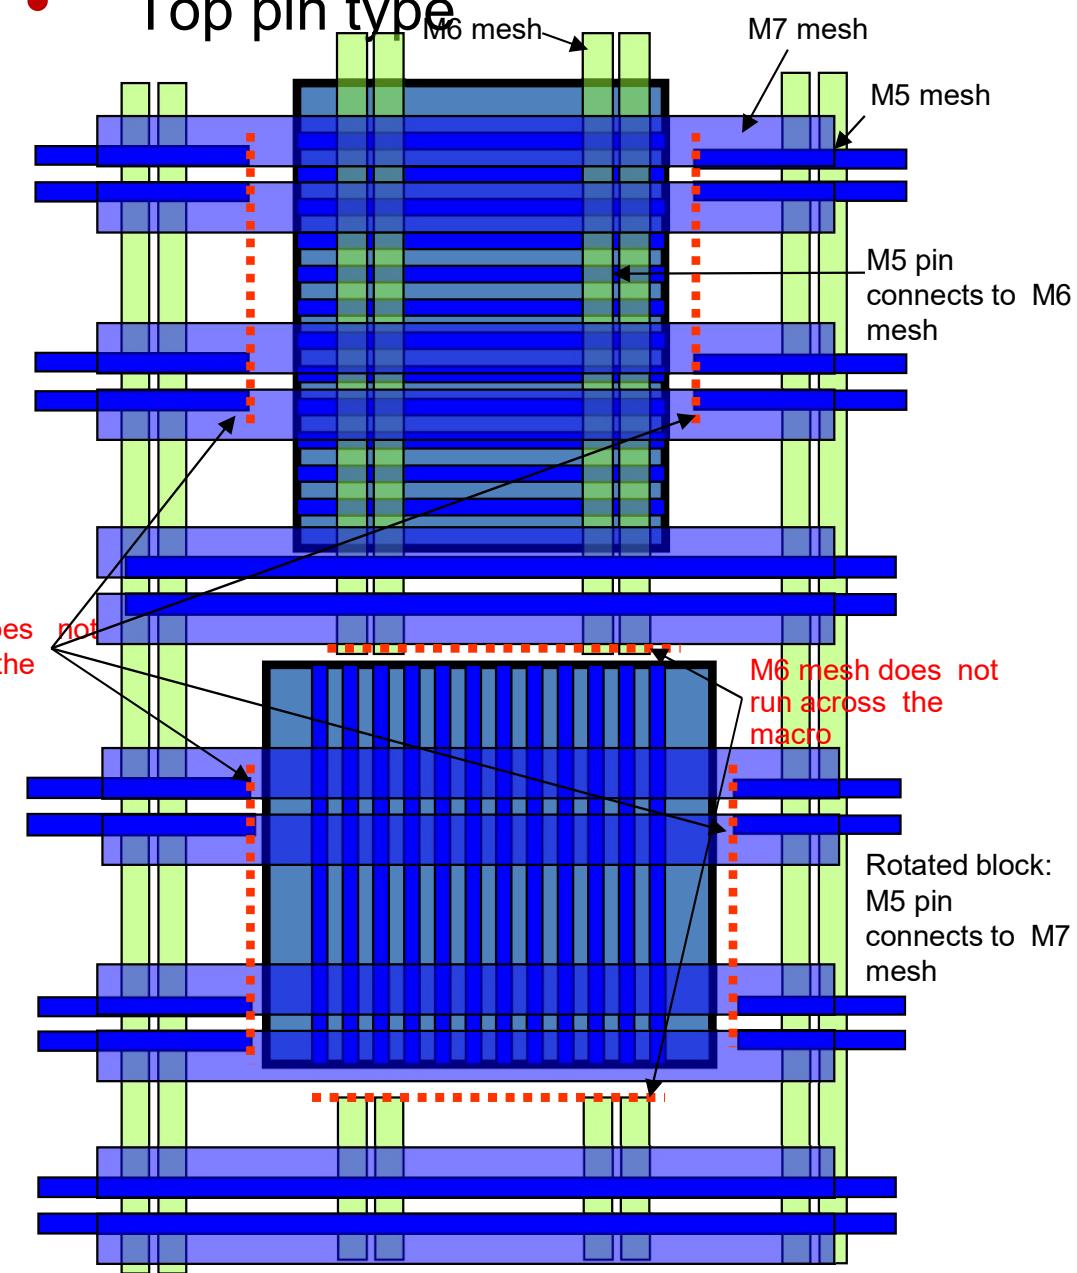


Macro Connection

- Side pin type



- Top pin type



pdn (aka *pdngen*) User Interface

- `set_voltage_domain` - define a new power domain
- `define_pdn_grid` - defines a new grid
- `add_pdn_stripe` - add straps and followpins
- `add_pdn_ring` - add rings
- `add_pdn_connect` - add layer connectivity
- `pdngen` - create grids

Main PDN Build Steps

1. Extract existing routing and blockages

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/PdnGen.cc#L112>

2. For each grid:

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/PdnGen.cc#L136>

a. Create rings, if requested

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c545b771cf5a7a9c1f4c71e869899ac978ffb5f/src/pdn/src/rings.cpp#L223>

b. Create straps

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c545b771cf5a7a9c1f4c71e869899ac978ffb5f/src/pdn/src/straps.cpp#L170>

c. Check connectivity and do channel repair (slide)

3. Estimate vias

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/PdnGen.cc#L148>

4. Trim excess

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/PdnGen.cc#L151>

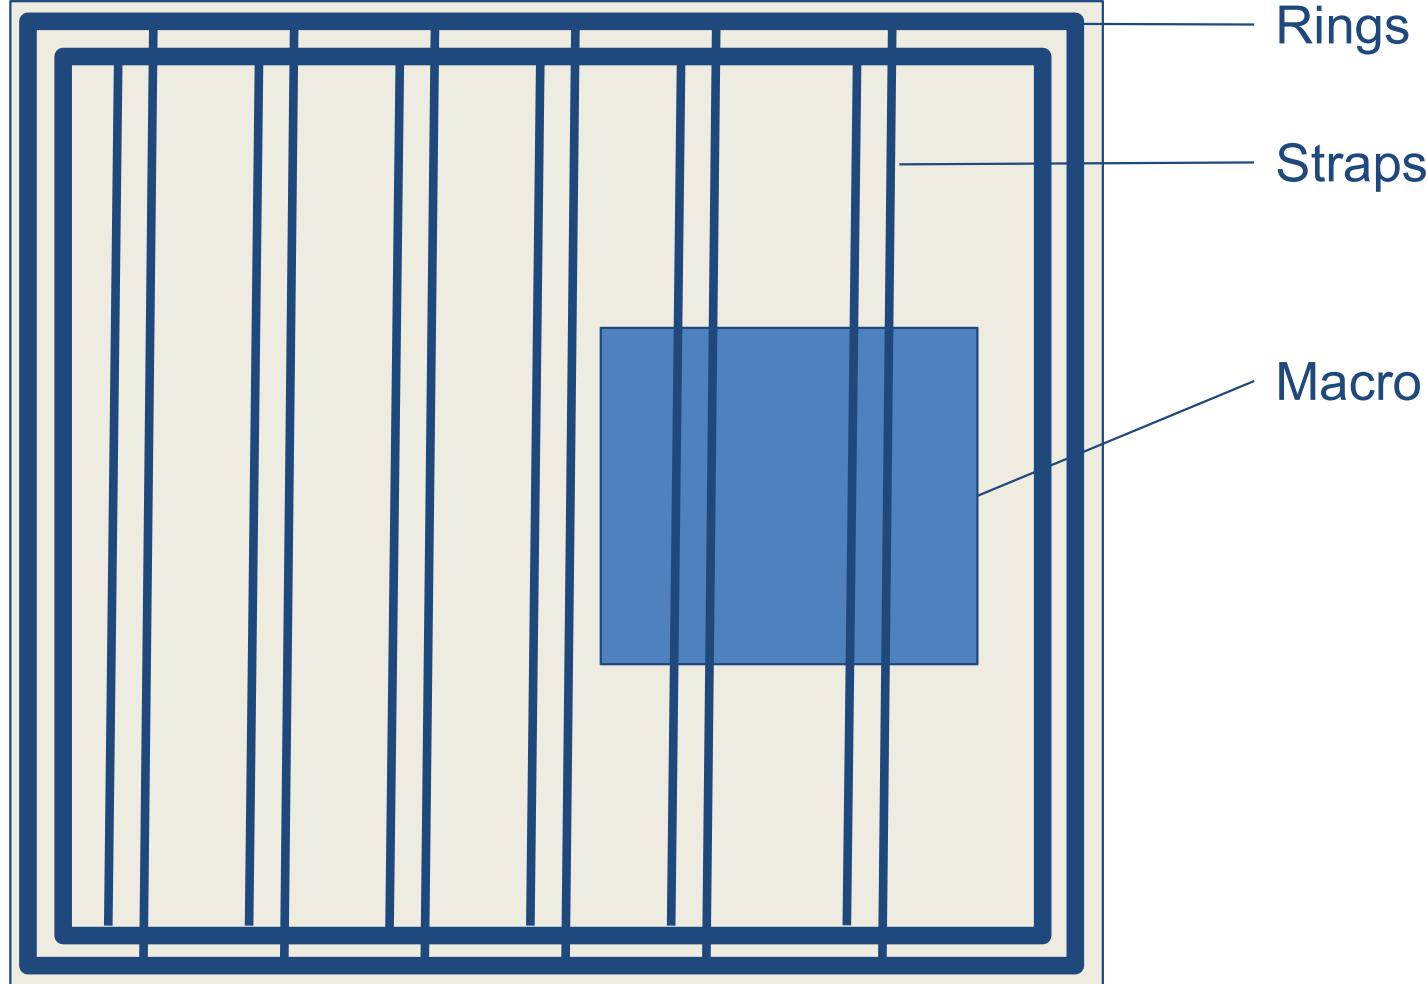
5. Commit to database

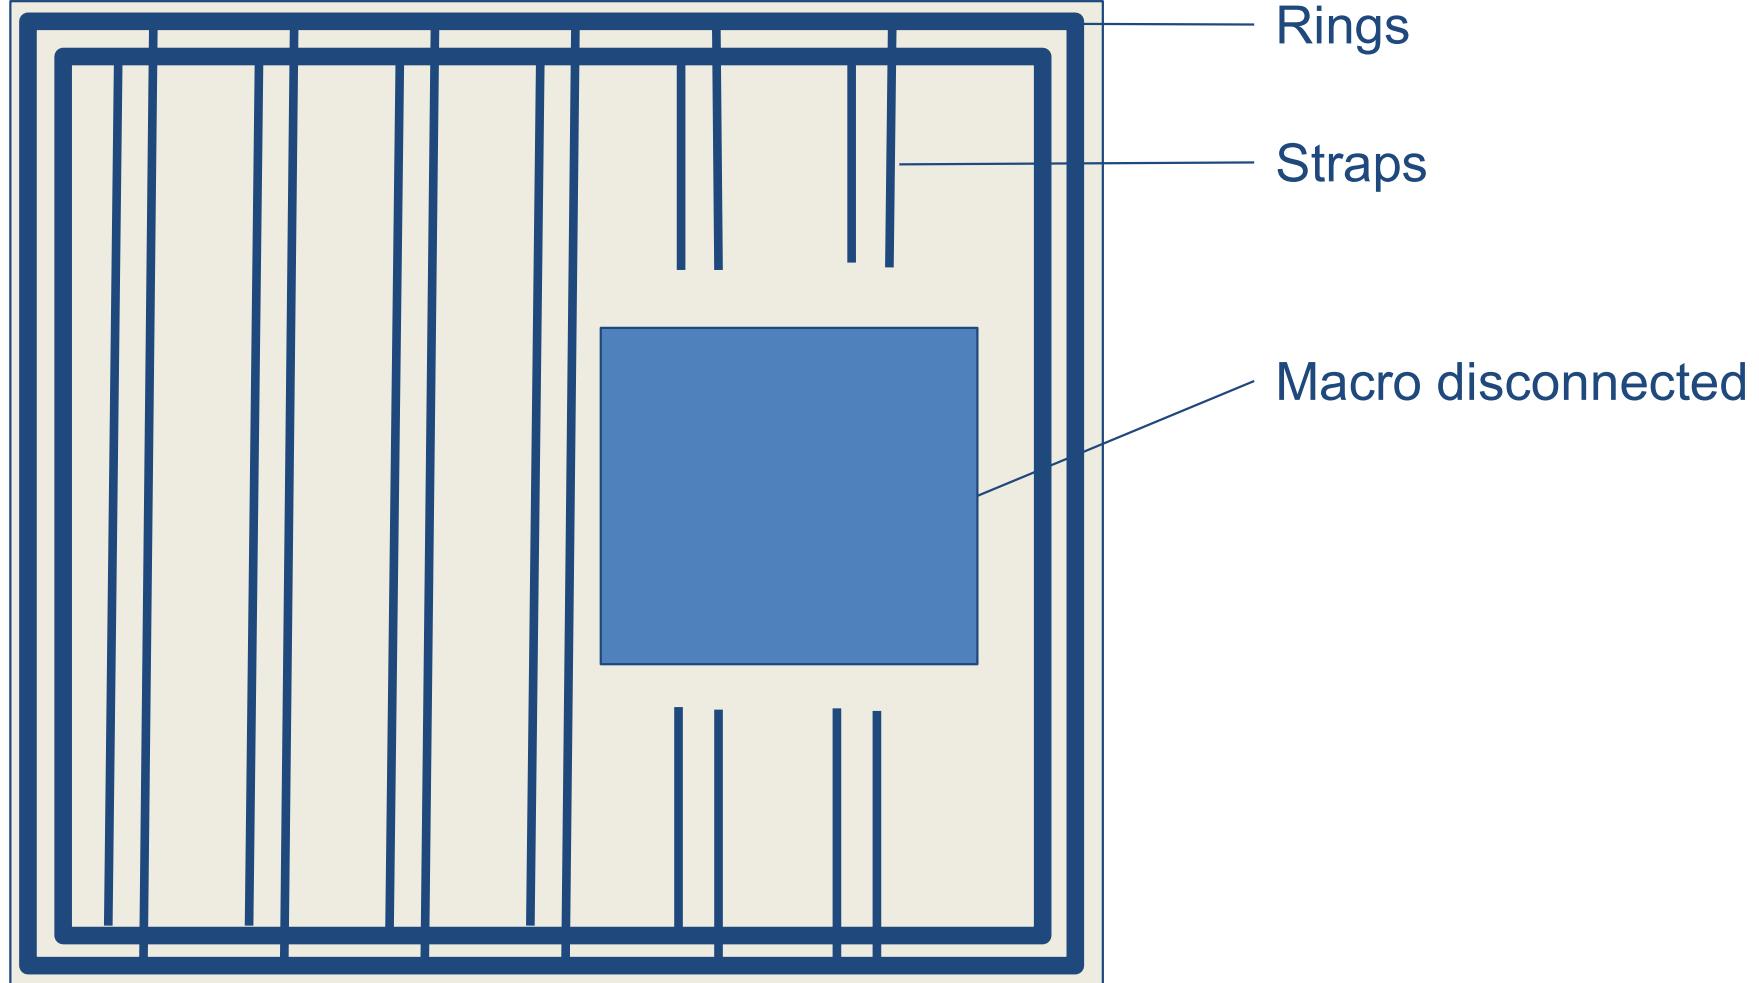
<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/PdnGen.cc#L731>

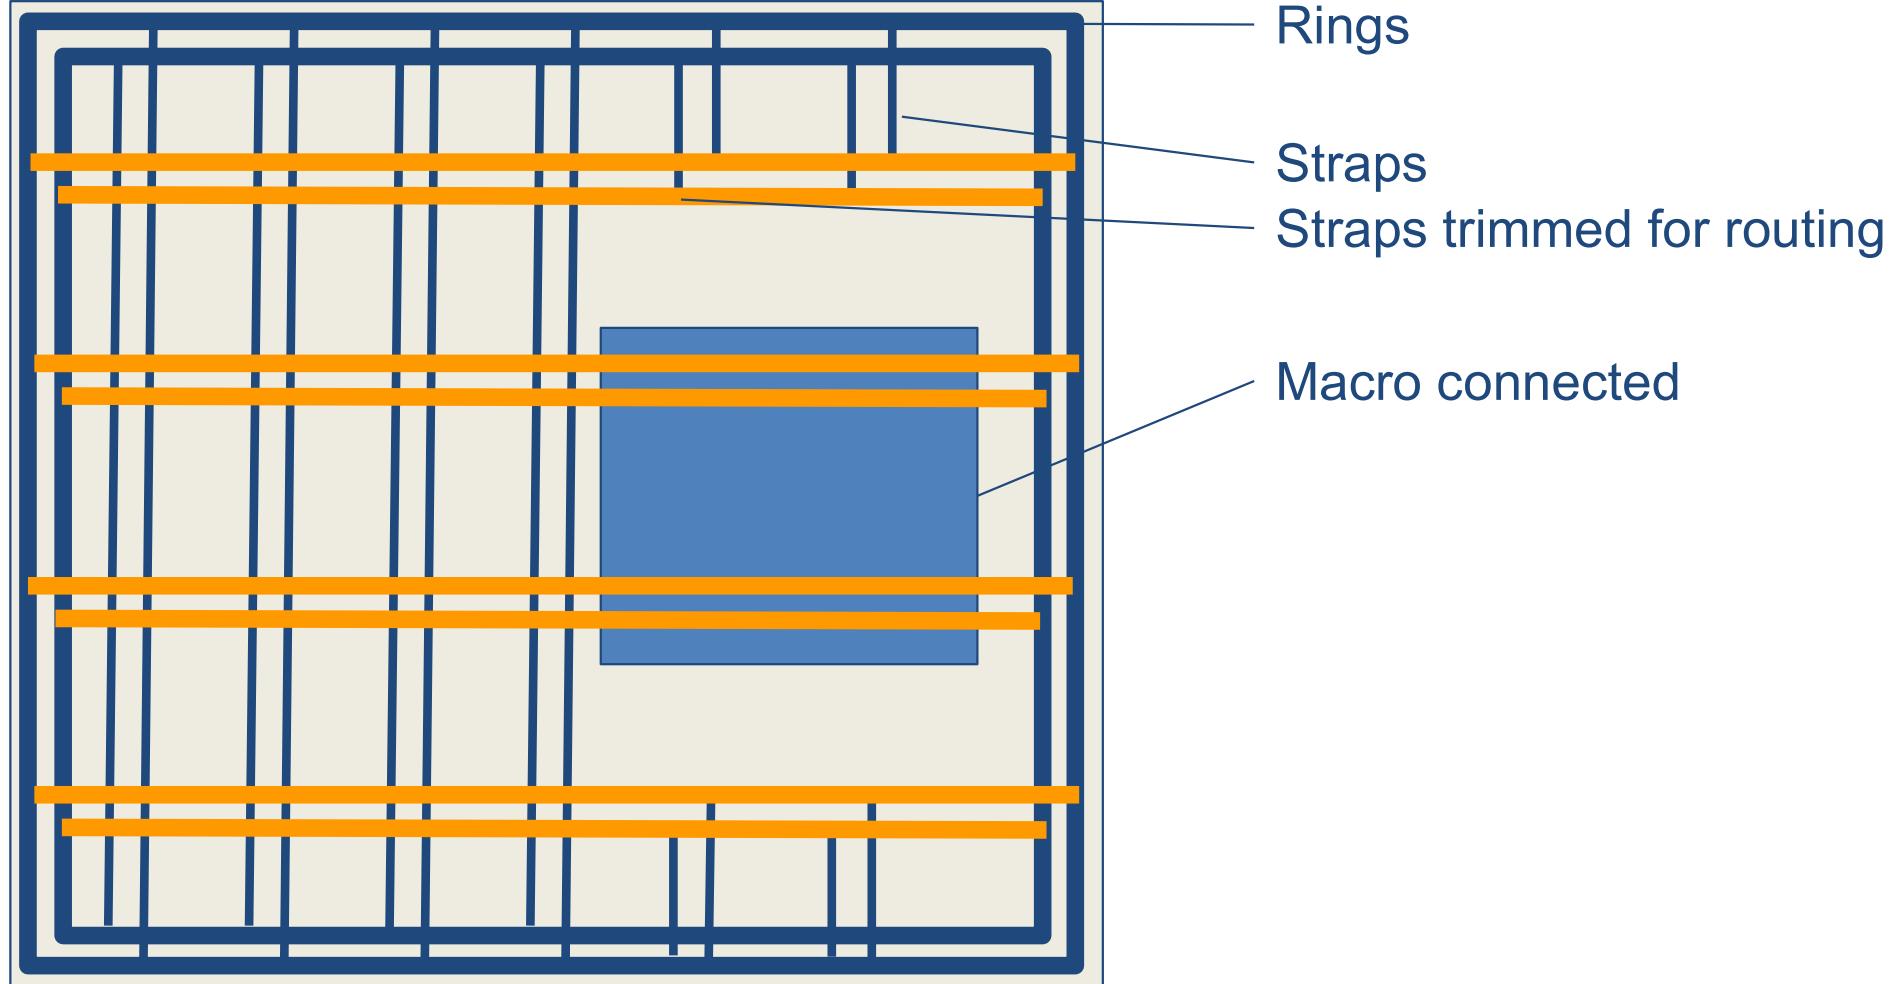
a. Create all vias DRC free

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/via.cpp>

b. Write wires to database







Channel Repair

<https://github.com/The-OpenROAD-Project/OpenROAD/blob/c34a30e0045eb68ad30c98f653e57f9b7f12d83c/src/pdn/src/straps.cpp#L1532>

1. Find disconnected channels (floating power grid shapes)
2. For each channel:
 - a. Find appropriate connection layers
 - b. Attempt repair in the center of the channels, if success, stop
 - c. Bisect channel and attempt repair in on left and right side of channel, if success, stop, else keep trying
 - i. Stop when bisection no longer can make progress

pdngen Improvements / Problem Statements

- Automatic connectivity rules (add_pdn_connect)
 - Create rules automatically for defining grid connectivity based on layers used, instances present, etc.
- Edge connectivity
 - Enable the power grid to connect to edge ports for macros, right now this is only possible on pad cells and standard cells, for macros we only connect from the top which requires macros use lower metals for power and routing. This is not a huge problem in nodes with a large number of metal layers, but for sky130, ihp130, etc, this imposes a large penalty.
- Power grid reinforcement (eco power grid)
 - Later in most flows, after detailed placement, it may be possible to determine if additional wires are needed to ensure IR drop stays within limits based on the power requirements of those areas. On the opposite side, it maybe possible to prune the grid if IR drop is not an issue and free up routing resources. [cf. <https://arxiv.org/abs/2110.14184>]
- Automatic power grid definition (hard)
 - Given an IR drop goal, create a power grid that meets this requirement based on the parasitic resistances and estimated power (either from placement or good guesses)

Paper References

- H. Murata, K. Fujiyoshi, S. Nakatake and Y. Kajitani, “VLSI module placement based on rectangle-packing by the sequence-pair”, IEEE Trans. on CAD, 15(12), 1996, pp. 1518-1524.
- X. Tang and D. F. Wong, “FAST-SP: a fast algorithm for block placement based on sequence pair”, Proc. ASP-DAC, 2001, pp. 521-526.
- S. N. Adya and I. L. Markov, “Fixed-outline Floorplanning : Enabling Hierarchical Design”, IEEE Trans. on VLSI Systems, 11(6), 2003, pp. 1120-1135.
- S. Kirkpatrick, C. D. Gelatt, Jr. and M. P. Vecchi, “Optimization by Simulated Annealing”, Science 220(4598) 1983, pp. 671-680.

BACKUP

B*-Tree Floorplan Representation

- Ordered binary tree (cf. “Packing Tree”, “O-Tree”)
- Root represents the block on the left-bottom corner
- Left child of node n_i represents the lowest unvisited block that belongs to the set of blocks located on the right-hand side of and adjacent to b_i

$$x_j = x_i + w_i$$

- Right child of the node n_i represents the lowest block located above and with x-coordinate equal to that of b_i

$$x_j = x_i$$

