

ECE 260C, Spring 2026

EDA Heuristic Optimizer's Toolkit

Andrew B. Kahng


About Me

Speaker Bio: Andrew B. Kahng is Distinguished Professor of CSE and ECE and holds the endowed chair in high-performance computing at UCSD. He has been a visiting scientist at Cadence and founder/CTO at Blaze DFM. He is coauthor of 3 books and over 500 journal and conference papers, holds 36 issued U.S. patents, and is a fellow of ACM and IEEE. He was the 2019 Ho-Am Prize laureate in Engineering. From 2000-2016 he served as international chair/co-chair of the International Technology Roadmap for Semiconductors (ITRS) Design and System Drivers working groups. He was PI of the DARPA “OpenROAD” project from June 2018 to December 2023, and was the founding PI and director of “TILOS” (<https://tilos.ai/>), an NSF AI Research Institute.

“EDA Heuristic Optimizer”

- **EDA**
 - Power, performance, area/cost, reliability, security == all about optimization
 - Semiconductor design automation: high stakes, ‘grungy’ (devices, wires, physics ...)
- **Heuristic** (as opposed to “algorithm”)
 - Algorithm = “recipe executed by/on a computer”
 - Heuristic = fewer guarantees
 - (and there are “metaheuristics” as well)
- **Optimizer**
 - Do better
 - Within resource limits


Design Optimization Lives in a Box





- **Start to End**: expensive!
 - $O(\text{year})$ for product
 - $O(\text{weeks})$ for SP&R and Opt
- Goal: best possible **End**
- Constraint: stay in “Box”
 - {compute}
 - X {licenses}
 - X {people}
 - X {weeks}

Huge space of trajectories: architecture, enablement, IPs, tools, manual fix, ...



RTL-to-GDS Chip Implementation Flow



IO Placement (Example: foundry 12nm RISC core, “coyote”)



Macro Placement



Tapcell (Well Tap) Insertion

Verilog
+ libraries,
constraints

Logic Synthesis

Floorplanning


Placement

Clock & Optimization



Global and Detailed
Routing

Layout Finishing



GDSII
final
layout





Power Delivery Network





Power Delivery Network – Zoom-In





Global Placement





Global Placement





Global Placement w/Partially Expanded Hierarchy





Global Placement Zoom-In





Sizing and Buffering (electrical rules)





(Legalized) Detailed Placement Zoom-In



Clock Tree Synthesis





Hold Fix Buffers





Congestion Map


Is this final placement routable?




Final (Detailed) Routing




What's Different (and Difficult) ?




- 1. Changing **abstractions**, **formats**, and **the design itself**



What's Different (and Difficult) ?



1. Changing **abstractions**, formats, and **the design itself**



2. Long **chains** of distinct, intractable **discrete optimizations**




- “Practical optimization” = metaheuristics on top of metaheuristics

Scale, multimodality, dynamism, diversity

→ 1000s of hidden commands and options in a commercial placer !





- Objectives are **ad hoc**
- Trajectories are **chaotic**
- Outcomes have **distributions**

Outcomes Are Heavily Mediated by (Heuristic) Optimizers !




- “ibm01” from ISPD-2002, “[Bookshelf](#)” format
- 12.7K instances
- FengShui ~2005; NTUplace3 ~2008; RePIAce ~2018


- “CT-Ariane133 X4”, [protobuf](#) format [\[link\]](#)
- 532 macros, 332K cells in TSMC 7nm
- AlphaChip 2024, SA/GWTW 1983/1994; RePIAce 2020-



What's Different (and Difficult) ?



1. Changing **abstractions**, formats, and **the design itself**



2. Long **chains** of distinct, intractable **discrete optimizations**

- “Practical optimization” = metaheuristics on top of metaheuristics


Scale, multimodality, dynamism, diversity

→ 1000s of hidden commands and options in a commercial placer !

- Objectives are **ad hoc**
- Trajectories are **chaotic**
- Outcomes have **distributions**

3. **Loops are expensive** (often, fatally so)

- Design process must “converge” both spatial embedding **and** performance




How?

With Chained Chaotic Optimizations ?!


Variation of metrics: $100 \times \left(\frac{\max}{\min} - 1 \right)$

3.4% Variation of Effective Clock Period




Logic
Synth

11% Variation of Total Cell Area




15% Variation of Effective Clock Period



Place
& Route

9% Variation of Total Cell Area



Implications for Optimization in IC Design

- Predictions today are **Constructive**


- Quick-and-dirty, under the hood

- Optimizations today are **Iterative**

- “Construct by Correction”


- Need Fast Simulations** to guide Optimization

- “Construct by Corrections ...
... that are **Correct by Construction**”



Catechism:

- A. You need **models** to have predictions
- B. You need predictions to **leverage** in exploration
- C. **What you can't predict, you guardband**
- D. **What you don't explore, you leave on the table**



ECE 260C, Spring 2026

EDA Heuristic Optimizer's Toolkit

Andrew B. Kahng

Mindsets

- **Theory vs. practice**
 - Big-O analysis
 - Performance ratios
 - ...
 - (benchmarks)
 - (real chips)
- **Does the problem matter?**
 - How do you know it matters (“skin in the game” etc.)?
 - Life cycles of (end-to-end) (collaborative) R&D
 - (Same effort to solve a relevant problem as an irrelevant problem ...)
 - Motivations, usage scenarios, validations (data, bandwidth, baselines, ...)
- **Delivery of value**
 - 10% or 10X
 - Post-RouteOpt PPA (rWL, TNS, FEP, power), not HPWL

Mindsets

- **Adversarial thinking**

- Can you defeat your heuristic? Make it perform badly? How?
- → “eyecharts”, “pathological examples”, “families/generators of instances” ...

- **Geometric thinking**

- 2D or mildly 3D (“2.5D”) domain

- **Context**

- Design and layout methodologies coevolve with tools and designs
 - Multi-core CPUs
 - SRAM-dominated SOC
 - Multimedia NOCs
 - Power and clock gating
 - DFT, BIST
 - ...
- Related: What is your path/strategy to demonstrating added value?



Heuristic (and Algorithmic) “Paradigms” = part of the toolkit

- **Divide and conquer**

- Small end cases
 - End-case partitioners and placers (“SmallPP” in Capo;
- Hypergraph partitioning and “VLSI-relevant” hypergraph partitioning
 - Hagen-Kahng spectral partitioning (ICCAD91), MLPart (ASPDAC00), SpecPart (ICCAD22), TritonPart (ICCAD23), FPGAPart, ChipletPart, ..., FastPart ...
 - Other side of the coin: clustering and sparsification; cf. multilevel, multiscale, multigrid ...
 - make the problem smaller (prevent mistakes, retain structure)

- **Greed**

- Goes a long way, especially with “multistart”
- And “orchestrations” (can pay small multiple of walltime)
 - Adaptive multistart (~evolutionary optimization of start points)
 - Go-With-The-Winners
 - (LSMC: large-step Markov chains ~”iterated descent”)



Heuristic (and Algorithmic) “Paradigms” = part of the toolkit

- **Dynamic programming**
 - Bottom-up, combine subsolutions
 - “what is your table and in what order do you fill it in?”
 - Need a sensible **ordering** of subproblems: left-to-right, bottom-up, topological order?, row-by-row?, ...
 - Van Ginneken optimal buffering/sizing
 - 1D has a natural order; 2D does not
 - Spacefilling curve used for CTS clustering – see “DP-RP” ([DAC-1994](#))
- **Geometric algorithms**
 - Manhattan geometry !
 - What is the maximum degree of a vertex in a minimum spanning tree? (In L_2 Euclidean plane? In L_1 Manhattan plane?)
 - How might you exploit the duality of L_∞ and L_1 ?

Heuristic (and Algorithmic) “Paradigms” = part of the toolkit

- **Combinatorial algorithms**
 - Network flows, matching
- **Mathematical programming**
 - Integer programming
 - Linear programming
 - + rounding
- **Satisfiability (SAT, SMT)**
- **Other**
 - Constraint programming?

Metaheuristics = also part of the toolkit

- **Metaphors**

- “explore vs. exploit”, “scatter vs. gather”, ...
- “greed vs. hill-climb vs. random walk”
- Physical analogies: electrostatics, masses on springs, ...

- **One thread**

- SA [TA, OBA, ...]
- LSMC = large-step Markov chain
- Tabu search

- **Populations of threads**

- (parallel multi-start = sampling; “grid search”, ...)
- GA
- Bayesian Optimization
- GWTW = go-with-the-winners

Does OpenROAD Use These? (of course ...)


- **Greed**
- **Dynamic Programming**
- **Simulated Annealing**
- **Multilevel Partitioning**
- **Physical Analogy**
- **Matching**
- **MILP**

1. Bayesian Optimization (BO)

- **Problem formulation:**
 - Given a **black-box** objective function $f(x)$ that is **expensive to evaluate**
 - Find $x^* = \arg \max f(x)$
- **Core idea:**
 - Build a **surrogate model (Gaussian Process)** that defines a distribution over the objective function
 - Use an **acquisition function** to decide where to sample next
 - Refine the surrogate model once new data is sampled


Source: [Link](#) [Link](#) [Link](#)

Illustration of Bayesian Optimization Process



Source: [Link](#) [Link](#) [Link](#)


Taxonomy of Bayesian Optimization Applications



Some popular implementations:
Ray Tune, Hyperopt, Vizier, ...


Source: [Link](#) [Link](#) [Link](#)

Example 1: SoC Architecture Design Exploration




Source: [Link](#)

Example 2: Physical Design Tool Parameter Tuning



Source: [Link](#)

Example: Cadence Cerebrus




Source: [Link](#)


Acknowledgment and thanks to Cadence Design Systems.

OpenROAD Flow Parameter Autotuning

- Automatic, iterative tuning to improve user-defined **score** within a given hyperparameter range space
- Interface: Python packages Ray/Tune
- Algorithms: HyperOpt, PBT, Optuna, Nevergrad, Ax, random search
- Advantages:
 - No need for pre-existing big data
 - Fewer trials needed than with, e.g., grid search
 - Parallelization management: core/thread mgmt, external server usage, visualization
 - Good outcomes within reasonable schedule, compute budgets



Flow Parameter AutoTuner – Architecture



Hyperparameter Space: SkyWater 130HD, ibex

- Tech: Skywater 130nm HD
- Design: ibex
- Tested Algorithms
 - HyperOpt
- Hyperparameter config <name, type, minmax, step>


```
1 {  
2   "GP_PAD": {"type": "int", "minmax": [0,4], "step": 1 },  
3   "DP_PAD": {"type": "int", "minmax": [0,4], "step": 1 },  
4   "LAYER_ADJUST": {"type": "float", "minmax": [0.1,0.7], "step": 0 },  
5   "PLACE_DENSITY": {"type": "float", "minmax": [0.1,1.0], "step": 0 },  
6   "FLATTEN": {"type": "int", "minmax": [0,1], "step": 1 },  
7   "PINS_DISTANCE": {"type": "int", "minmax": [1,3], "step": 1 },  
8   "CTS_CLUSTER_SIZE": {"type": "int", "minmax": [10,40], "step": 1 },  
9   "CTS_CLUSTER_DIAMETER": {"type": "int", "minmax": [80,120], "step": 1 },  
10  "GR_OVERFLOW": {"type": "int", "minmax": [1,1], "step": 0 }  
11 }
```

Assuming full factorial combinations,
1,058,298,150
= #possible combinations!

- When type is int and step = 0, it means constant value
- When type is float and step = 0, it means continuous range


TensorBoard Visualization: SkyWater 130HD, ibex

- GUI integration with **TensorBoard**
- Score results versus Wall Time



“Interactive Learning”

- **AI that explores, interacts with, and learns in real-world environments**
- **Bandits**
 - Multi-arm (no state information)
 - Contextual (state information independent of past actions)
- **Bayesian Optimization** (global optimization = infinite bandits)
- **Active Learning**
- **Reinforcement Learning** (state information dependent on past actions)




Source: [Link](#)
See also: [Link](#)

BACKUP / SPARE (Lecture 9)

Multi-Layer Routing on a Grid Graph

- 2D routes are assigned layers. Then, refinement is done over the 3D graph
- Example of the first 4 layers as a grid graph

Courtesy of, and copyright by, Dr. Patrick Groeneveld, Stanford EE 292
Please do not redistribute without Dr. Groeneveld's permission.




Layer M4:
Vertical preference

Layer M3:
Horizontal preference

Layer M2:
Vertical preference


Layer M1:
Horizontal preference

Vertical edge
represents a via




Grid Graph Capacity Model in GRT

- GRT abstracts detailed geometry into GCells
- Each edge carries capacity, usage, and overflow
 - $\sum \text{overflow}$ is the primary routability objective.
- Capacity based on:
 - Number of tracks and layers,
 - Track preferred directions,
 - Blockages,
 - And user-adjustable derates.
- 2D-first approach
 - Per-layer 3D resources are summed into the 2D accounting used by the primary global-routing loop




(a)




(b)

Pattern and Monotonic Routing

- Global routing scales by using fast initial pattern routing
 - L/Z shapes, three-bend patterns
- Monotonic routing provides a deterministic set of routes
- Loop of: select pattern, test congestion/blockages, select best
 - Then, apply repair (ripup-and-reroute)




L-Shaped



Z-Shaped


Pattern Routing



Monotonic Routing

Maze Routing

- Point to point routing of nets
 - from source to sink
- Basic idea = wave propagation (Lee, 1961)
 - Breadth-first search + back-tracing after finding shortest path
 - Guarantees to find the shortest path
- Objective = route all nets according to some cost function that minimizes congestion, route length, coupling, etc.




Cost = label
= distance from "A"

Finding Connections: Dijkstra vs. A*

Courtesy of, and copyright by, Dr. Patrick Groeneveld, Stanford EE 292
Please do not redistribute without Dr. Groeneveld's permission.

Dijkstra (“BFS with Alarm Clocks”) vs. A* (“Best-First Search”)





Source: Wikipedia


- Even with obstacles present, given an optimistic lower bound $h(n)$ on “future cost”, A* (aka “best-first search”) will guarantee to find a lowest-cost path from Source to Target
- Speed on a grid graph with $n \times n$ nodes: $O(n)$ best case, $O(n^2)$ worst case

Incremental Ripup and Reroute


Courtesy of, and copyright by, Dr. Patrick Groeneveld, Stanford EE 292
Please do not redistribute without Dr. Groeneveld's permission.



Rip up offending
nets



Route Net2





Then route Net1

Note: this was
simply a lucky
guess

- Small “switchbox” (in DRT)
 - Detailed routing problem is typically decomposed into switchboxes, with solutions then stitched back together ...
 - GRT problem is over (sets of) Steiner tree (edges)
- Routed sequentially, net by net (Net1 before Net2 → spacing violation)
- In GRT, we see congestion/routability violations that force ripup-reroute

Connecting Multi-Terminal Nets


- In general, maze and line-probe routing are not well-suited to multi-terminal nets
- Several attempts made to extend to multi-terminal nets
 - Connect one terminal at a time
 - Use the entire connected subtrees as sources or targets during expansion
- Improve solution quality
 - Ripup/Reroute to improve solution quality (remove a segment and re-connect)
 - In GR: some topological perturbation (edge shifting)



- Results are sub-optimal
- Inherit time and memory cost of maze and line-probe algorithms

Topology Selection

- Multi-pin nets must be decomposed before two-pin routing: topology is the first routing decision for each net.
- For a given pin set: the chosen tree affects the balance of congestion, wirelength, via count, and timing
- OpenROAD's `stt` implements fast FLUTE RSMTs for minimizing wirelength.
 - Routability and timing need more




Changing topology moves routing demand before maze routing starts, even though the pin set is unchanged.

Prim-Dijkstra Construction

- Timing-aware trees trade off source-sink radius against total wirelength.


But large delay from source for nodes 3,4,5




Prim's Minimum Spanning Tree (MST)
Minimizes wirelength (WL)

Dijkstra's Shortest Path Tree (SPT)
Minimizes source-sink pathlengths (PLs)

But large tree wirelength!



Prim-Dijkstra (PD) tradeoff of Prim's MST and Dijkstra's SPT



Directly trades off the Prim, Dijkstra constructions

Prim-Dijkstra Gives a Timing-Aware Tree Knob

- MST-like trees minimize total wirelength; SPT-like trees minimize source-to-sink pathlength.
- Prim-Dijkstra blends the two with α , letting topology trade routing cost against timing radius.
- This matters because minimum-wirelength trees can create large detours, while shortest-path trees can waste routing resources.
- The idea is simple enough to implement greedily and useful as intuition for timing-aware topology selection.

Prim-Dijkstra core

Choose edge e_{ij} minimizing
 $d_{ij} + \alpha \cdot l_i$

$\alpha = 0$ -> Prim-like low-WL tree

$\alpha = 1$ -> Dijkstra-like low-radius tree

Greedy update:

add (i,j) to T
set $l_j = l_i + d_{ij}$

<https://vlscad.ucsd.edu/Publications/Journals/j18.pdf>

<https://vlscad.ucsd.edu/Publications/Conferences/355/c355.pdf>

Capacity Control: Global Margin and Local Steering

- Layer adjustment globally scales the capacity available to GRT, creating detailed-routing margin
- Macro extension reserves a halo region around macros
- Region adjustment changes capacity inside a chosen rectangle to steer demand away from identified potential hotspots

Capacity controls

Global margin:

```
set_global_routing_layer_adjustment ...
```

Macro halo:

```
set_macro_extension 2
```

Local hotspot steering:

```
set_global_routing_region_adjustment ...
```

Derate Example: ASAP7

- ORFS ASAP7 enablement deliberately scales effective routing capacity **down by 25%**.
 - This spreads demand earlier and helps avoid detailed routing getting stuck on overly tight guides.
- Pessimism acknowledges advanced-node pin access, via limits, and DRC effects that the global graph cannot model effectively.

```
ASAP7 fastroute.tcl

set_global_routing_layer_adjustment \
    $::env(MIN_ROUTING_LAYER) -
    $::env(MAX_ROUTING_LAYER) 0.25

set_routing_layers -clock \
    $::env(MIN_CLK_ROUTING_LAYER) -
    $::env(MAX_ROUTING_LAYER)

set_routing_layers -signal \
    $::env(MIN_ROUTING_LAYER) -
    $::env(MAX_ROUTING_LAYER)
```


Challenges for Global Router Development

- Runtime and memory scalability
- Congestion prediction
 - Congestion-aware global routing?
 - History-based cost functions (NCR)?
- Net routing order
 - How to maximize the concurrency of routing decisions?
 - ILP or multi-commodity flow based methods?
 - Rip-up and reroute (or Reroute w/o Ripup)?
- Steiner tree topology reconstruction
 - Steiner point relocation
 - Complete solution space exploration
 - RSMT or RMST?
- Timing, placement-synthesis services, fidelity to detailed routing, ...

Tool users have opportunities to guide the global router – weights, priorities, non-default rules, layer ranges, density screens, track supply reductions, ...
See, e.g., <https://vlsicad.ucsd.edu/Publications/Conferences/411/c411.pdf>

GRT Commands and Knobs

- `set_routing_layers` and layer adjustment define where traffic may go and how much capacity GRT should believe.
- `set_macro_extension` and region-level adjustments steer traffic away from local hotspots before DRT fails.
- `-congestion_iterations`, `-critical_nets_percentage`, and `-skip_large_fanout_nets` trade quality, timing focus, and runtime.
- `-allow_congestion`, incremental reroute flags, and experimental hooks expose triage or specialized behavior.

```
global_route -guide_file route.guide \  
  -congestion_iterations 50 \  
  -critical_nets_percentage 30 \  
  -congestion_report_file congestion.rpt  
  
set_routing_layers -signal M2-M10 -clock M6-M9  
set_global_routing_layer_adjustment M4-M8 0.30  
set_macro_extension 2
```

When GRT Fails, Match the Symptom to the Knob

- Global overflow everywhere usually means placement is too dense: reduce utilization or add whitespace before over-tuning the router.
- Layer-specific or macro-edge hotspots point to layer adjustment, layer ranges, macro extension, or region adjustment.
- If guides look clean but DRT still fails, investigate pin access or local geometric issues rather than just global congestion.
 - `draw_route_segments {my_net}`
- Near-converged runs may need more congestion iterations; noisy high-fanout nets may need to be skipped for debug.

CTS

Clock Distribution

- General goal of clock distribution
 - Deliver clock to all memory elements with acceptable skew
 - Deliver clock edges with acceptable sharpness
- Clocking network design is one of the greatest challenges in the design of a large chip
- Clocks generally distributed via wiring trees (and meshes)
- Low-resistance interconnect to minimize delay, variability
- Low-capacitance interconnect to minimize power
- Multiple drivers to distribute clock signal
 - Use optimal sizing principles to design buffers
 - Clock lines can create significant crosstalk

Issues in Clock Distribution Network Design



- Skew
 - Process, voltage, and temperature (especially across “wide corners” – 1.25V to 0.70V, 125C to -40C, etc.) <https://vlsicad.ucsd.edu/Publications/Conferences/327/c327.pdf>
 - Data dependence
 - Noise coupling
 - Load balancing
- Power, CV^2f
 - Clock gating
- Flexibility/Tunability
 - Compactness – fit into existing layout/design
- Reliability
 - Electromigration

OpenROAD CTS Commands

Command	Description	Example Output
<code>clock_tree_synthesis</code>	Build a balanced Htree by choosing appropriate clock buffers	<pre> [INFO CTS-0050] Root buffer is BUF_X4. [INFO CTS-0051] Sink buffer is BUF_X4. [INFO CTS-0052] The following clock buffers will be used for CTS: BUF_X4 ... [INFO CTS-0017] Max level of the clock tree: 5. [INFO CTS-0098] Clock net "clk" [INFO CTS-0099] Sinks 2537 [INFO CTS-0100] Leaf buffers 96 [INFO CTS-0101] Average sink wire length 9247.25 um [INFO CTS-0102] Path depth 18 - 19 [INFO CTS-0207] Leaf load cells 62 [INFO RSZ-0058] Using max wire length 693um. [INFO RSZ-0047] Found 33 long wires. [INFO RSZ-0048] Inserted 94 buffers in 33 nets. </pre>
<code>configure_cts_characterization</code>	<pre>-max_slew <ps> -max_cap <fF> -slew_steps <N> -cap_steps <N></pre> Set maximum values for characterization to test against. Set number of discrete buckets "steps" that slew/cap values will be divided into.	N/A
<code>repair_clock_nets</code>	Fixes LDRC violations including max wire length	<pre>[INFO RSZ-0058] Using max wire length 2154um.</pre>
<code>report_clock_skew</code>	Report worst clock skew for each clock signal in the design	<pre> Clock clk 1.26 source latency inst_7_12/clk ^ -1.13 target latency inst_8_12/clk ^ 0.00 CRPR ----- 0.13 setup skew </pre>
<code>report_checks -format full_clock_expanded</code>	Report timing violations including clock paths	<pre> Startpoint: dp.rf.rf[31][3]\$_DFFE_PP_ (rising edge-triggered Flip-flop clocked by clk) Endpoint: aluout[0] (output port clocked by clk) Path Group: clk Path Type: max Fanout Cap Slew Delay Time Description ----- 0.00 0.00 clock clk (rise edge) 0.00 0.00 clock source latency 1 0.09 0.00 0.00 0.00 ^ clk (in) 0.00 0.00 0.00 ^ clk (net) 0.00 0.00 0.00 ^ clkbuf_0_clk/A (sky130_fd_sc_hd__clkbuf_16) 8 0.21 0.22 0.25 0.25 ^ clkbuf_0_clk/X (sky130_fd_sc_hd__clkbuf_16) </pre>

CTS Main Steps

- Sink clustering
 - Sequential elements are grouped into a fixed number of clusters based on their locations
- Tree construction and balancing
 - Buffers are inserted based on some structure, e.g., hierarchical H-Tree
 - Tree lengths are balanced such that clock skews are minimized
- LDRC (“ERC”, “DRV”) **electrical** rules repair
 - LDRC violations are repaired during or after CTS
 - Max transition, max capacitance, max wire length, etc.



Sink Clustering

- Group sequential elements based on their locations to produce the best results (e.g., minimum wire length)
 - These parameters can be specified manually or determined automatically
 - Cluster size
 - Cluster diameter
- All elements in the cluster will be driven by the same buffer





Figure 2: The planar Sierpinski spacefilling curve.


<https://vlsicad.ucsd.edu/Publications/Conferences/32/c32.pdf>

Early “TritonCTS” versions used spacefilling curves to perform sink clustering !

Sink Clustering Implementation

- `normalizePoints` maps sinks into a unit box
- `computeTheta` recursively computes a locality-preserving order key θ
- `sortPoints` orders sinks by theta
- `findBestMatching` tests groupings and tracks best cost
- `isLimitExceeded` enforces size, diameter, or cap-based limits

```
normalizePoints(maxDiameter);  
computeAllThetas();  
sortPoints();  
bestSolutionFound =  
    findBestMatching(groupSize);  
  
limit:    size >= groupSize  
         or diameter > maxInternalDiameter  
         or cap > sinkBufferInputCap * factor
```




Sink Clustering Flow

- Collect sink coordinates and insertion-delay data
- Run SinkClustering over candidate size/diameter combinations
 - Choose the best clustering solution found
- Create leaf buffers such as `clkbuf_leaf_*`
- Create leaf clock nets and attach clustered sinks

Obstruction-Aware CTS

- Clock buffers should not be placed on top of macros, placement blockages or another clock buffers
- Detailed placement may displace “illegal” buffers and cause timing to change after CTS
- New “legal” buffer locations need to preserve balanced clock tree
- Obstruction-aware CTS can reduce legalizer displacement by up to 4X



sky130hd/microwatt **without**
obstruction-aware CTS




sky130hd/microwatt **with**
obstruction-aware CTS


CTS Origin: Generalized H-Tree Concept

- Structured clock trees

K. Han, A. B. Kahng and J. Li, "Optimal Generalized H-Tree Topology and Buffering for High-Performance and Low-Power Clock Distribution", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* <https://vlsicad.ucsd.edu/Publications/Journals/j128.pdf>.



H-tree



"Fishbone"

	H-tree	Fishbone
Skew		<
Wirelength		>
Latency		>
Power		>



Generalized H-tree (GH-tree)


Can we mix two clock structures to have better tradeoff between clock power vs. skew or latency?

- History: (1) Bakoglu’s 1988 book made H-tree approach well-known. Cadence CTGen (Dr. Lars Hagen), mid-1990s, started trend toward “fishbone” style – save capacitance!
- Starting in the 2010s: on-chip variation (OCV) derates became very costly, so goal became to reduce insertion delay (== “latency”) along with power.


Generalized H-Tree Concept

- Structured clock trees

K. Han, A. B. Kahng and J. Li, "Optimal Generalized H-Tree Topology and Buffering for High-Performance and Low-Power Clock Distribution", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* <https://vlscad.ucsd.edu/Publications/Journals/j128.pdf>.




H-tree



"Fishbone"

	H-tree	Fishbone
Skew		<
Wirelength		>
Latency		>
Power		>




Generalized H-tree (GH-tree)

a balanced tree topology
with arbitrary branching
factor at each level

GH-tree with depth $P = 8$ and branching factors (4, 2, 2, 2, 4, 2, 2, 2)


Idea: Capture/Explore Tradeoff (“Pareto” Frontier)

- (Recall: floorplan shape functions ?)
- GH-tree explores branching patterns across levels
- Power, skew, latency, and wirelength create a Pareto tradeoff
- Buffering and topology must be co-optimized




Embedding of GH-Tree Into a Sink Placement


- Clustering and branching point / buffer placement using Integer Linear Programming (ILP) and Linear Programming (LP)



Layout Example of GH-tree on VGA




GH-tree (brown) with the downstream clock trees (black) within sink regions



GH-tree (brown) with clustered downstream sinks


Buffer Selection and Electrical Limits

- Root and sink (LCB) buffer choices finalized during characterization setup
- Fanout limits are read from SDC and Liberty where available
- Sink clustering sizes can be limited by LCB (leaf clock buffer) fanout limits
- Resizer max-wire-length guidance can be used after CTS
- Derates adjust pessimism for sink buffer max cap and delay buffers




Clock Roots, Gated Clocks, and Builders

- Clock roots found from OpenSTA clock definitions and database nets
- Each valid clock tree gets an HTreeBuilder
- Gated clocks create hierarchical or separate clock-tree construction contexts
- Top input nets and parent builders preserve hierarchy relationships
- Recent changes: avoid recursion through clock-gating loops, etc.



Hierarchy and Macro/Register Balancing

- Macro clock pins may have existing internal insertion delay
- Register sinks and macro sinks should have balanced arrival times
- Generated/gated clock trees can form nested builder relationships
- Latency balancing is applied after tree construction/writeback context exists



NDR, Repair, and Electrical Cleanup

- `-apply_ndr` controls whether clock nets get non-default-rule treatment
- `-repair_clock_nets` invokes electrical repair for long/high-fanout clock nets
- `-distance_between_buffers` controls wire segmentation
- `-sink_buffer_max_cap_derate` and `-delay_buffer_derate` tune pessimism
- Post-CTS repair is crucial for final QOR (fmax)

NDR mode	Meaning
none	no clock NDR application
root_only	apply near root only
half	partial clock tree NDR strategy
full	apply broadly through clock tree


Nondefault Rules (NDRs)

- NDRs apply wider wire widths and spacings that are different from the default rules specified in technology Library Exchange Format (LEF) files as long as they do not violate any design rules (DRC)
 - widening wires decreases sheet resistance $\Rightarrow$ smaller wire delay
 - NDR examples: double-width and –spacing (2W2S); triple-width and –spacing (3W3S), ...
- NDRs are used to fix electromigration (EM) violations, reduce parasitic and delay variations
 - e.g., in clock routing to meet insertion delay/skew requirements
 - e.g., at post-route stage to fix EM violations by widening violating net segments
- Side effects of widening wires with NDRs
 - Increase in driven capacitance $\Rightarrow$ increase in dynamic power


Use of NDRs in Clock Routing

- Eliminate noise
 - Coupling capacitance induces delay uncertainty. Larger spacings can reduce delay uncertainty.
- Prevent EM violation
 - When larger buffers drive minimum-width wires, they overshoot the current density limit. So, wider wires MUST be used to prevent EM violations
- Reduce variation
 - Smaller geometries and more resistive interconnects, in conjunction with multi-patterning in lithography ($< 20\text{nm}$), result in large parasitics and delay variations. Wider wires typically have less variation
- Reduce power (complementary: reduce skew)
 - Wide wires increase capacitance which leads to increase in dynamic power. Applying NDRs smartly (i.e. by wire-tapering), total capacitance (hence power) can be reduced

Specification of NDRs in LEF, DEF



Track Impact of NDRs





- Track cost = Number of routing tracks unavailable to the detailed router because of increase in the required spacing to neighboring wires when applying NDRs
- Key: Router centers each wire segment on a track gridline

After CTS Construction: Database Writeback

- CTS creates clock buffer instances and connects them to generated subnets.
- Uses stable prefixes: `clkbuf_*`, `clknet_*`, `clkload_*`, `delaybuf_*`
- Can create dummy loads unless disabled
- Applies NDR strategy when requested
- Counts and reports sinks after writing the tree


Clock Tree Viewer

- Open GUI
 - `gui::show`
- Enable “Clock Tree Viewer” if not enabled
- Clock tree viewer shows latencies at all sinks
 - Red sinks = FF/latches
 - Green sinks = macros
 - Insertion delays are added to macro sinks



Summary

- OpenROAD CTS is a database-integrated, based clock-tree builder
- Sink clustering and obstruction-aware legalization, “GH-tree” antecedents, “spacefilling-curve clustering” are part of CTS even today.
- Many enhancements over time: simplified characterization, NDR support, sink polarity and hierarchy support, etc.
 - = progression from GH-tree concepts to more robust hierarchical CTS



Watch These Spaces ...

- CTS
 - MBFFs (this year's MP1)
 - OCV
 - Clock-data co-opt
 - Useful skew and retiming
- GRT
 - Timing-driven (with budgeting!)
 - RC-aware (beyond 'per-dbu-per-layer' or 'per-via' costs)
 - Better parasitic estimations earlier
 - Empowered to change the placement
 - Blurring with the detailed router
 - GPU acceleration (no easy decomposition into subproblems as in GRT)